

# k-Nearest Neighbor (kNN) exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

The kNN classifier consists of two stages:

- During training, the classifier takes the training data and simply remembers it
- During testing, kNN classifies every test image by comparing to all training images and transferring the labels of the k most similar training examples
- The value of k is cross-validated

In this exercise you will implement these steps and understand the basic Image Classification pipeline, cross-validation, and gain proficiency in writing efficient, vectorized code.

In [1]:

```
# Run some setup code for this notebook.

import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

# This is a bit of magic to make matplotlib figures appear inline in the notebook
# rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

In [2]:

```
# Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

# Cleaning up variables to prevent loading data multiple times (which may cause
memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

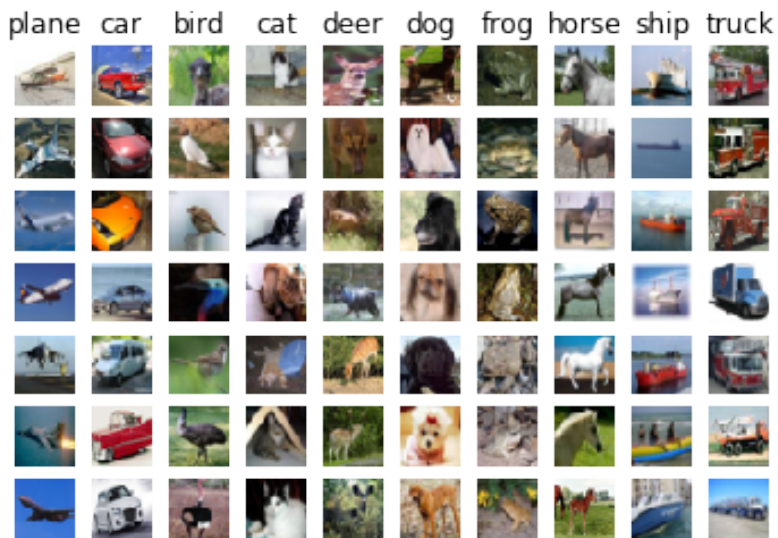
X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

In [3]:

```
# Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



In [4]:

```
# Subsample the data for more efficient code execution in this exercise
num_training = 5000
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
num_test = 500
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
```

In [5]:

```
# Reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
print(X_train.shape, X_test.shape)
```

```
(5000, 3072) (500, 3072)
```

In [6]:

```
from cs231n.classifiers import KNearestNeighbor
```

```
# Create a kNN classifier instance.
# Remember that training a kNN classifier is a noop:
# the Classifier simply remembers the data and does no further processing
classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
```

We would now like to classify the test data with the kNN classifier. Recall that we can break down this process into two steps:

1. First we must compute the distances between all test examples and all train examples.
2. Given these distances, for each test example we find the  $k$  nearest examples and have them vote for the label

Lets begin with computing the distance matrix between all training and test examples. For example, if there are  $N_{tr}$  training examples and  $N_{te}$  test examples, this stage should result in a  $N_{te} \times N_{tr}$  matrix where each element  $(i,j)$  is the distance between the  $i$ -th test and  $j$ -th train example.

First, open `cs231n/classifiers/k_nearest_neighbor.py` and implement the function `compute_distances_two_loops` that uses a (very inefficient) double loop over all pairs of (test, train) examples and computes the distance matrix one element at a time.

In [7]:


```
# Open cs231n/classifiers/k_nearest_neighbor.py and implement
# compute_distances_two_loops.
```

```
# Test your implementation:
dists = classifier.compute_distances_two_loops(X_test)
print(dists.shape)
```

```
(500, 5000)
```

In [8]:

```
# We can visualize the distance matrix: each row is a single test example and
# its distances to training examples
plt.imshow(dists, interpolation='none')
plt.show()
```



**Inline Question #1:** Notice the structured patterns in the distance matrix, where some rows or columns are visible brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

- What in the data is the cause behind the distinctly bright rows?
- What causes the columns?

**Your Answer:**

A distinctly bright row indicates that this test example is dissimilar (has large distance) to examples in the train set. A distinctly bright column indicates that this training example is dissimilar to test examples.

In [9]:

```
# Now implement the function predict_labels and run the code below:
# We use k = 1 (which is Nearest Neighbor).
y_test_pred = classifier.predict_labels(dists, k=1)

# Compute and print the fraction of correctly predicted examples
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 137 / 500 correct => accuracy: 0.274000

You should expect to see approximately 27% accuracy. Now let's try out a larger k, say k = 5:

In [10]:

```
y_test_pred = classifier.predict_labels(dists, k=5)
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))
```

Got 139 / 500 correct => accuracy: 0.278000

You should expect to see a slightly better performance than with k = 1.

**Inline Question 2** We can also other distance metrics such as L1 distance. The performance of a Nearest Neighbor classifier that uses L1 distance will not change if (Select all that apply.):

1. The data is preprocessed by subtracting the mean.
2. The data is preprocessed by subtracting the mean and dividing by the standard deviation.
3. The coordinate axes for the data are rotated.
4. None of the above.

Your Answer: 1

Your explanation: let  $x_1, x_2$  be the original vectors. L1 distance is then  $|x_1 - x_2|$  let  $m$  be the mean;  $|(x_1 - m) - (x_2 - m)| = |x_1 - x_2|$  let  $s$  be the std;  $|(x_1 - m)/s - (x_2 - m)/s| = |(x_1 - x_2)/s|$  let  $r$  be rotation matrix;  $x_{1\_new}$  and  $x_{2\_new}$  can then become perpendicular to  $x_1 - x_2$

In [11]:

```
# Now lets speed up distance matrix computation by using partial vectorization
# with one loop. Implement the function compute_distances_one_loop and run the
# code below:
dists_one = classifier.compute_distances_one_loop(X_test)

# To ensure that our vectorized implementation is correct, we make sure that it
# agrees with the naive implementation. There are many ways to decide whether
# two matrices are similar; one of the simplest is the Frobenius norm. In case
# you haven't seen it before, the Frobenius norm of two matrices is the square
# root of the squared sum of differences of all elements; in other words, reshap
e
# the matrices into vectors and compute the Euclidean distance between them.
difference = np.linalg.norm(dists - dists_one, ord='fro')
print('Difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

Difference was: 0.000000

Good! The distance matrices are the same

In [12]:

```
# Now implement the fully vectorized version inside compute_distances_no_loops
# and run the code
dists_two = classifier.compute_distances_no_loops(X_test)
```

In [13]:

```
# check that the distance matrix agrees with the one we computed before:
difference = np.linalg.norm(dists - dists_two, ord='fro')
print('Difference was: %f' % (difference, ))
if difference < 0.001:
    print('Good! The distance matrices are the same')
else:
    print('Uh-oh! The distance matrices are different')
```

Difference was: 0.000000

Good! The distance matrices are the same

In [14]:

```
# Let's compare how fast the implementations are
def time_function(f, *args):
    """
    Call a function f with args and return the time (in seconds) that it took to
    execute.
    """
    import time
    tic = time.time()
    f(*args)
    toc = time.time()
    return toc - tic

two_loop_time = time_function(classifier.compute_distances_two_loops, X_test)
print('Two loop version took %f seconds' % two_loop_time)

one_loop_time = time_function(classifier.compute_distances_one_loop, X_test)
print('One loop version took %f seconds' % one_loop_time)

no_loop_time = time_function(classifier.compute_distances_no_loops, X_test)
print('No loop version took %f seconds' % no_loop_time)

# you should see significantly faster performance with the fully vectorized impl
ementation
```

Two loop version took 34.598362 seconds

One loop version took 90.215767 seconds

No loop version took 0.241169 seconds

## Cross-validation

We have implemented the k-Nearest Neighbor classifier but we set the value  $k = 5$  arbitrarily. We will now determine the best value of this hyperparameter with cross-validation.

In [15]:

```
num_folds = 5
k_choices = [1, 3, 5, 8, 10, 12, 15, 20, 50, 100]
```

```
X_train_folds = []
y_train_folds = []
#####
# TODO:
# Split up the training data into folds. After splitting, X_train_folds and
# y_train_folds should each be lists of length num_folds, where
# y_train_folds[i] is the label vector for the points in X_train_folds[i].
# Hint: Look up the numpy array_split function.
#####
X_train_folds = np.array_split(X_train, num_folds)
y_train_folds = np.array_split(y_train, num_folds)
#####
#                               END OF YOUR CODE
#####

# A dictionary holding the accuracies for different values of k that we find
# when running cross-validation. After running cross-validation,
# k_to_accuracies[k] should be a list of length num_folds giving the different
# accuracy values that we found when using that value of k.
k_to_accuracies = {}

# X_train, y_train, X_test, y_test
#####
# TODO:
# Perform k-fold cross validation to find the best value of k. For each
# possible value of k, run the k-nearest-neighbor algorithm num_folds times,
# where in each case you use all but one of the folds as training data and the
# last fold as a validation set. Store the accuracies for all fold and all
# values of k in the k_to_accuracies dictionary.
#####
for i in range(num_folds):
    X_train_i = np.concatenate(X_train_folds[:i] + X_train_folds[i+1:])
    y_train_i = np.concatenate(y_train_folds[:i] + y_train_folds[i+1:])

    X_test_i = X_train_folds[i]
    y_test_i = y_train_folds[i]

    for k in k_choices:
        classifier.train(X_train_i, y_train_i)
        y_test_pred = classifier.predict(X_test_i, k)
        num_correct = np.sum(y_test_pred == y_test_i)
        accuracy = float(num_correct) / len(y_test_i)
        #print('Got %d / %d correct => accuracy: %f' % (num_correct, len(y_test_
i), accuracy))
        curr_dict = k_to_accuracies.get(k, [])
        curr_dict.append(accuracy)
        k_to_accuracies[k] = curr_dict

#####
#                               END OF YOUR CODE
#####
```



*# Print out the computed accuracies*


```
for k in sorted(k_to_accuracies):  
    for accuracy in k_to_accuracies[k]:  
        print('k = %d, accuracy = %f' % (k, accuracy))
```

k = 1, accuracy = 0.263000  
k = 1, accuracy = 0.257000  
k = 1, accuracy = 0.264000  
k = 1, accuracy = 0.278000  
k = 1, accuracy = 0.266000  
k = 3, accuracy = 0.239000  
k = 3, accuracy = 0.249000  
k = 3, accuracy = 0.240000  
k = 3, accuracy = 0.266000  
k = 3, accuracy = 0.254000  
k = 5, accuracy = 0.248000  
k = 5, accuracy = 0.266000  
k = 5, accuracy = 0.280000  
k = 5, accuracy = 0.292000  
k = 5, accuracy = 0.280000  
k = 8, accuracy = 0.262000  
k = 8, accuracy = 0.282000  
k = 8, accuracy = 0.273000  
k = 8, accuracy = 0.290000  
k = 8, accuracy = 0.273000  
k = 10, accuracy = 0.265000  
k = 10, accuracy = 0.296000  
k = 10, accuracy = 0.276000  
k = 10, accuracy = 0.284000  
k = 10, accuracy = 0.280000  
k = 12, accuracy = 0.260000  
k = 12, accuracy = 0.295000  
k = 12, accuracy = 0.279000  
k = 12, accuracy = 0.283000  
k = 12, accuracy = 0.280000  
k = 15, accuracy = 0.252000  
k = 15, accuracy = 0.289000  
k = 15, accuracy = 0.278000  
k = 15, accuracy = 0.282000  
k = 15, accuracy = 0.274000  
k = 20, accuracy = 0.270000  
k = 20, accuracy = 0.279000  
k = 20, accuracy = 0.279000  
k = 20, accuracy = 0.282000  
k = 20, accuracy = 0.285000  
k = 50, accuracy = 0.271000  
k = 50, accuracy = 0.288000  
k = 50, accuracy = 0.278000  
k = 50, accuracy = 0.269000  
k = 50, accuracy = 0.266000  
k = 100, accuracy = 0.256000  
k = 100, accuracy = 0.270000  
k = 100, accuracy = 0.263000  
k = 100, accuracy = 0.256000  
k = 100, accuracy = 0.263000

In [16]:

```
# plot the raw observations
for k in k_choices:
    accuracies = k_to_accuracies[k]
    plt.scatter([k] * len(accuracies), accuracies)

# plot the trend line with error bars that correspond to standard deviation
accuracies_mean = np.array([np.mean(v) for k,v in sorted(k_to_accuracies.items())])
accuracies_std = np.array([np.std(v) for k,v in sorted(k_to_accuracies.items())])
plt.errorbar(k_choices, accuracies_mean, yerr=accuracies_std)
plt.title('Cross-validation on k')
plt.xlabel('k')
plt.ylabel('Cross-validation accuracy')
plt.show()
```



In [18]:

```
# Based on the cross-validation results above, choose the best value for k,
# retrain the classifier using all the training data, and test it on the test
# data. You should be able to get above 28% accuracy on the test data.
best_k = 6

classifier = KNearestNeighbor()
classifier.train(X_train, y_train)
y_test_pred = classifier.predict(X_test, k=best_k)

# Compute and display the accuracy
num_correct = np.sum(y_test_pred == y_test)
accuracy = float(num_correct) / num_test
print('Got %d / %d correct => accuracy: %f' % (num_correct, num_test, accuracy))

Got 141 / 500 correct => accuracy: 0.282000
```

**Inline Question 3** Which of the following statements about  $k$ -Nearest Neighbor ( $k$ -NN) are true in a classification setting, and for all  $k$ ? Select all that apply.

1. The training error of a 1-NN will always be better than that of 5-NN.
2. The test error of a 1-NN will always be better than that of a 5-NN.
3. The decision boundary of the  $k$ -NN classifier is linear.
4. The time needed to classify a test example with the  $k$ -NN classifier grows with the size of the training set.
5. None of the above.

*Your Answer:* 1, 4

*Your explanation:* Training error of 1-NN is 0; the closest neighbor is always itself. For testing on the other hand, we introduce new examples that may be labeled better by considering more data. Since at test time KNN must calculate distance to all training examples, time to classify test grows with train set size. KNN decision boundaries are arbitrary depending on distribution of data, most often not linear.

# Multiclass Support Vector Machine exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

In this exercise you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

In [1]:

```
# Run some setup code for this notebook.

import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

# This is a bit of magic to make matplotlib figures appear inline in the
# notebook rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

## CIFAR-10 Data Loading and Preprocessing

In [2]:

```
# Load the raw CIFAR-10 data.
cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

# Cleaning up variables to prevent loading data multiple times (which may cause
memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

Training data shape: (50000, 32, 32, 3)

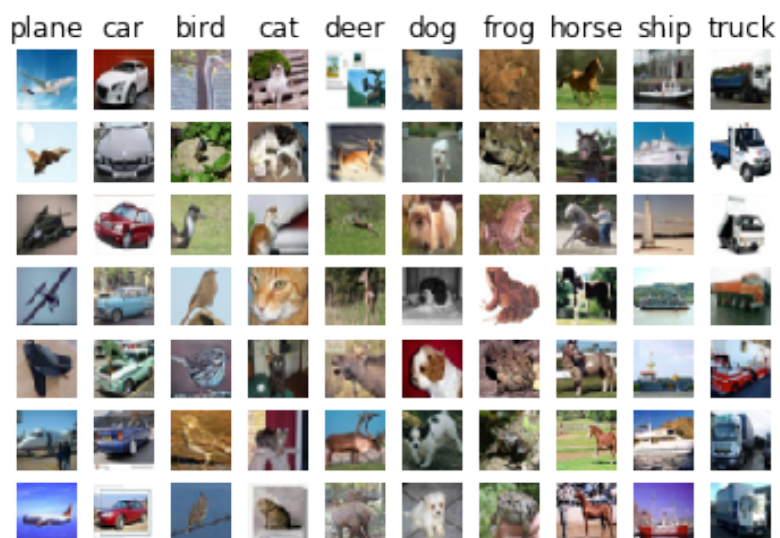
Training labels shape: (50000,)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000,)

In [3]:

```
# Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



In [4]:

```
# Split the data into train, val, and test sets. In addition we will
# create a small development set as a subset of the training data;
# we can use this for development so our code runs faster.
num_training = 49000
num_validation = 1000
num_test = 1000
num_dev = 500

# Our validation set will be num_validation points from the original
# training set.
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]

# Our training set will be the first num_train points from the original
# training set.
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

# We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# We use the first num_test points of the original test set as our
# test set.
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
```



In [5]:

```
# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))


# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('dev data shape: ', X_dev.shape)
```

```
Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

In [6]:

```
# Preprocessing: subtract the mean image
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)
print(mean_image[:10]) # print a few of the elements
plt.figure(figsize=(4,4))
plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean i
mage
plt.show()
```

```
[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
 131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]
```



In [7]:

```
# second: subtract the mean image from train and test data
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image
```

In [8]:

```
# third: append the bias dimension of ones (i.e. bias trick) so that our SVM
# only has to worry about optimizing a single weight matrix W.
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)

(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)
```

## SVM Classifier

Your code for this section will all be written inside **cs231n/classifiers/linear\_svm.py**.

As you can see, we have prefilled the function `compute_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

In [9]:

```
# Evaluate the naive implementation of the loss we provided for you:
from cs231n.classifiers.linear_svm import svm_loss_naive
import time

# generate a random SVM weight matrix of small numbers
W = np.random.randn(3073, 10) * 0.0001

loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
print('loss: %f' % (loss, ))
```

loss: 8.828588

The grad returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you:

In [10]:

```
# Once you've implemented the gradient, recompute it with the code below
# and gradient check it with the function we provided for you

# Compute the loss and its gradient at W.
loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)

# Numerically compute the gradient along several randomly chosen dimensions, and
# compare them with your analytically computed gradient. The numbers should match
# almost exactly along all dimensions.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad)
print("****")

# do the gradient check once again with regularization turned on
# you didn't forget the regularization gradient did you?
loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad)
```

```
numerical: 8.951727 analytic: 8.951727, relative error: 4.156012e-11
numerical: 20.910217 analytic: 20.863606, relative error: 1.115787e-03
numerical: 30.920956 analytic: 30.970534, relative error: 8.010515e-04
numerical: 9.897412 analytic: 9.897412, relative error: 6.348008e-12
numerical: 17.382492 analytic: 17.382492, relative error: 2.776074e-11
numerical: -7.964577 analytic: -8.001215, relative error: 2.294777e-03
numerical: 10.839305 analytic: 10.839305, relative error: 2.566910e-11
numerical: -23.446868 analytic: -23.446868, relative error: 7.532532e-12
numerical: -11.284455 analytic: -11.284455, relative error: 1.187367e-11
numerical: 19.946680 analytic: 19.946680, relative error: 1.960008e-11
****
numerical: -4.493948 analytic: -4.524749, relative error: 3.415241e-03
numerical: 7.383087 analytic: 7.383087, relative error: 2.509571e-11
numerical: 17.474259 analytic: 17.474259, relative error: 2.633592e-12
numerical: -4.830305 analytic: -4.830305, relative error: 8.046636e-11
numerical: -7.765563 analytic: -7.765563, relative error: 6.169870e-11
numerical: 6.656028 analytic: 6.656028, relative error: 3.426538e-11
numerical: 24.989061 analytic: 24.952233, relative error: 7.374411e-04
numerical: -9.231988 analytic: -9.243166, relative error: 6.050103e-04
numerical: 21.204751 analytic: 21.224611, relative error: 4.680770e-04
numerical: 0.307938 analytic: 0.307938, relative error: 9.428135e-10
```

## Inline Question 1:

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? How would change the margin affect of the frequency of this happening? *Hint: the SVM loss function is not strictly speaking differentiable*

**Your Answer:** Discrepancy can be caused by an incorrect numerical gradient due to the loss function being indifferentiable and hence taking the difference across such a boundary is invalid, such as at 0. Changing the margin would change how often / where we switch from no loss to loss; the effect on frequency is unclear (depends on how often the difference between class scores hits this new boundary).

In [11]:

```
# Next implement the function svm_loss_vectorized; for now only compute the loss
;
# we will implement the gradient in a moment.
tic = time.time()
loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))

from cs231n.classifiers.linear_svm import svm_loss_vectorized
tic = time.time()
loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# The losses should match but your vectorized implementation should be much faster.
print('difference: %f' % (loss_naive - loss_vectorized))
```

Naive loss: 8.828588e+00 computed in 0.085618s  
Vectorized loss: 8.828588e+00 computed in 0.027589s  
difference: -0.000000

In [12]:

```
# Complete the implementation of svm_loss_vectorized, and compute the gradient
# of the loss function in a vectorized way.

# The naive implementation and the vectorized implementation should match, but
# the vectorized version should still be much faster.
tic = time.time()
_, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Naive loss and gradient: computed in %fs' % (toc - tic))

tic = time.time()
_, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Vectorized loss and gradient: computed in %fs' % (toc - tic))

# The loss is a single number, so it is easy to compare the values computed
# by the two implementations. The gradient on the other hand is a matrix, so
# we use the Frobenius norm to compare them.
difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('difference: %f' % difference)
```

Naive loss and gradient: computed in 0.081041s  
Vectorized loss and gradient: computed in 0.011492s  
difference: 0.000000

## Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss.


In [13]:

```
# In the file linear_classifier.py, implement SGD in the function  
# LinearClassifier.train() and then run it with the code below.  
from cs231n.classifiers import LinearSVM  
svm = LinearSVM()  
tic = time.time()  
loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,  
                      num_iters=1500, verbose=True)  
toc = time.time()  
print('That took %fs' % (toc - tic))
```

```
iteration 0 / 1500: loss 789.930274  
iteration 100 / 1500: loss 286.482087  
iteration 200 / 1500: loss 108.001109  
iteration 300 / 1500: loss 42.434364  
iteration 400 / 1500: loss 18.525659  
iteration 500 / 1500: loss 10.049786  
iteration 600 / 1500: loss 7.151333  
iteration 700 / 1500: loss 5.557743  
iteration 800 / 1500: loss 5.563516  
iteration 900 / 1500: loss 5.457811  
iteration 1000 / 1500: loss 5.429021  
iteration 1100 / 1500: loss 5.470428  
iteration 1200 / 1500: loss 5.677682  
iteration 1300 / 1500: loss 5.349342  
iteration 1400 / 1500: loss 5.057903  
That took 8.804816s
```

In [14]:

```
# A useful debugging strategy is to plot the loss as a function of
# iteration number:
plt.plot(loss_hist)
plt.xlabel('Iteration number')
plt.ylabel('Loss value')
plt.show()
```



In [15]:

```
# Write the LinearSVM.predict function and evaluate the performance on both the
# training and validation set
y_train_pred = svm.predict(X_train)
print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
y_val_pred = svm.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
```

```
training accuracy: 0.371061
validation accuracy: 0.380000
```

In [19]:

```
# Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of about 0.4 on the validation set.
learning_rates = [1e-7, 2e-7, 4e-7]
regularization_strengths = [1e4, 2e4, 3e4]

# results is dictionary mapping tuples of the form
# (learning_rate, regularization_strength) to tuples of the form
# (training_accuracy, validation_accuracy). The accuracy is simply the fraction
# of data points that are correctly classified.
results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation rate
```

```
best_params = None
```

```
#####  
# TODO: #  
# Write code that chooses the best hyperparameters by tuning on the validation #  
# set. For each combination of hyperparameters, train a linear SVM on the #  
# training set, compute its accuracy on the training and validation sets, and #  
# store these numbers in the results dictionary. In addition, store the best #  
# validation accuracy in best_val and the LinearSVM object that achieves this #  
# accuracy in best_svm. #  
# #  
# Hint: You should use a small value for num_iters as you develop your #  
# validation code so that the SVMs don't take much time to train; once you are #  
# confident that your validation code works, you should rerun the validation #  
# code with a larger value for num_iters. #  
#####  
for l in learning_rates:  
    for r in regularization_strengths:  
        svm = LinearSVM()  
        curr_loss = svm.train(X_train, y_train, learning_rate=l, reg=r,  
                               num_iters=1500, verbose=True)  
        y_train_pred = svm.predict(X_train)  
        y_train_acc = np.mean(y_train == y_train_pred)  
        y_val_pred = svm.predict(X_val)  
        y_val_acc = np.mean(y_val == y_val_pred)  
        results[(l, r)] = (y_train_acc, y_val_acc)  
        if y_val_acc > best_val:  
            best_svm = svm  
            best_val = y_val_acc  
            best_params = (l, r)  
#####  
#                                     END OF YOUR CODE #  
#####  
  
# Print out results.  
for lr, reg in sorted(results):  
    train_accuracy, val_accuracy = results[(lr, reg)]  
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (  
        lr, reg, train_accuracy, val_accuracy))  
  
print('best validation accuracy achieved during cross-validation: %f' % best_val  
)  
print(best_params)
```

```
iteration 0 / 1500: loss 329.812881  
iteration 100 / 1500: loss 213.884048  
iteration 200 / 1500: loss 143.058440  
iteration 300 / 1500: loss 96.680817  
iteration 400 / 1500: loss 66.274272  
iteration 500 / 1500: loss 45.652302  
iteration 600 / 1500: loss 32.018373  
iteration 700 / 1500: loss 22.933892
```



iteration 800 / 1500: loss 16.960706  
iteration 900 / 1500: loss 13.204388  
iteration 1000 / 1500: loss 9.905237  
iteration 1100 / 1500: loss 8.624346  
iteration 1200 / 1500: loss 7.805575  
iteration 1300 / 1500: loss 6.477198  
iteration 1400 / 1500: loss 5.921364  
iteration 0 / 1500: loss 641.142059  
iteration 100 / 1500: loss 284.166304  
iteration 200 / 1500: loss 129.317181  
iteration 300 / 1500: loss 60.885383  
iteration 400 / 1500: loss 29.764385  
iteration 500 / 1500: loss 16.033179  
iteration 600 / 1500: loss 10.002331  
iteration 700 / 1500: loss 7.221273  
iteration 800 / 1500: loss 6.113470  
iteration 900 / 1500: loss 5.898611  
iteration 1000 / 1500: loss 5.647968  
iteration 1100 / 1500: loss 4.748945  
iteration 1200 / 1500: loss 5.250188  
iteration 1300 / 1500: loss 5.457525  
iteration 1400 / 1500: loss 5.639028  
iteration 0 / 1500: loss 939.636591  
iteration 100 / 1500: loss 281.393841  
iteration 200 / 1500: loss 87.032507  
iteration 300 / 1500: loss 30.018033  
iteration 400 / 1500: loss 12.058377  
iteration 500 / 1500: loss 7.707717  
iteration 600 / 1500: loss 6.032632  
iteration 700 / 1500: loss 5.942680  
iteration 800 / 1500: loss 6.080024  
iteration 900 / 1500: loss 5.455763  
iteration 1000 / 1500: loss 5.577070  
iteration 1100 / 1500: loss 5.593244  
iteration 1200 / 1500: loss 5.278550  
iteration 1300 / 1500: loss 5.370906  
iteration 1400 / 1500: loss 5.448102  
iteration 0 / 1500: loss 334.880441  
iteration 100 / 1500: loss 145.749133  
iteration 200 / 1500: loss 67.019249  
iteration 300 / 1500: loss 31.935991  
iteration 400 / 1500: loss 16.874266  
iteration 500 / 1500: loss 10.518698  
iteration 600 / 1500: loss 7.668947  
iteration 700 / 1500: loss 6.304107  
iteration 800 / 1500: loss 6.116000  
iteration 900 / 1500: loss 5.828414  
iteration 1000 / 1500: loss 5.485808  
iteration 1100 / 1500: loss 4.864327  
iteration 1200 / 1500: loss 5.101638  
iteration 1300 / 1500: loss 5.231340  
iteration 1400 / 1500: loss 5.233736  
iteration 0 / 1500: loss 647.722236

iteration 100 / 1500: loss 129.030187  
iteration 200 / 1500: loss 29.647306  
iteration 300 / 1500: loss 10.381281  
iteration 400 / 1500: loss 6.229066  
iteration 500 / 1500: loss 5.287628  
iteration 600 / 1500: loss 5.034977  
iteration 700 / 1500: loss 5.086883  
iteration 800 / 1500: loss 4.900851  
iteration 900 / 1500: loss 5.071654  
iteration 1000 / 1500: loss 4.995439  
iteration 1100 / 1500: loss 5.349900  
iteration 1200 / 1500: loss 4.650375  
iteration 1300 / 1500: loss 5.032852  
iteration 1400 / 1500: loss 6.045765  
iteration 0 / 1500: loss 943.513361  
iteration 100 / 1500: loss 87.700810  
iteration 200 / 1500: loss 12.691915  
iteration 300 / 1500: loss 6.496296  
iteration 400 / 1500: loss 5.773099  
iteration 500 / 1500: loss 5.794557  
iteration 600 / 1500: loss 6.278239  
iteration 700 / 1500: loss 5.329029  
iteration 800 / 1500: loss 5.808075  
iteration 900 / 1500: loss 5.530021  
iteration 1000 / 1500: loss 5.997961  
iteration 1100 / 1500: loss 5.569494  
iteration 1200 / 1500: loss 5.711300  
iteration 1300 / 1500: loss 5.374030  
iteration 1400 / 1500: loss 5.113521  
iteration 0 / 1500: loss 333.361183  
iteration 100 / 1500: loss 67.944055  
iteration 200 / 1500: loss 17.117683  
iteration 300 / 1500: loss 7.442095  
iteration 400 / 1500: loss 5.785281  
iteration 500 / 1500: loss 5.392824  
iteration 600 / 1500: loss 5.244573  
iteration 700 / 1500: loss 4.841607  
iteration 800 / 1500: loss 5.774334  
iteration 900 / 1500: loss 5.524201  
iteration 1000 / 1500: loss 5.364770  
iteration 1100 / 1500: loss 4.612733  
iteration 1200 / 1500: loss 5.402273  
iteration 1300 / 1500: loss 5.708605  
iteration 1400 / 1500: loss 4.692609  
iteration 0 / 1500: loss 637.681800  
iteration 100 / 1500: loss 28.715797  
iteration 200 / 1500: loss 6.398879  
iteration 300 / 1500: loss 5.323034  
iteration 400 / 1500: loss 5.697103  
iteration 500 / 1500: loss 5.497378  
iteration 600 / 1500: loss 4.910578  
iteration 700 / 1500: loss 5.173264  
iteration 800 / 1500: loss 6.031431


```
iteration 900 / 1500: loss 5.393203
iteration 1000 / 1500: loss 5.647651
iteration 1100 / 1500: loss 5.567541
iteration 1200 / 1500: loss 5.107683
iteration 1300 / 1500: loss 5.211513
iteration 1400 / 1500: loss 5.565286
iteration 0 / 1500: loss 951.727452
iteration 100 / 1500: loss 12.247992
iteration 200 / 1500: loss 5.634766
iteration 300 / 1500: loss 6.143236
iteration 400 / 1500: loss 5.552660
iteration 500 / 1500: loss 5.745416
iteration 600 / 1500: loss 5.421951
iteration 700 / 1500: loss 5.998420
iteration 800 / 1500: loss 5.485592
iteration 900 / 1500: loss 5.410751
iteration 1000 / 1500: loss 5.983546
iteration 1100 / 1500: loss 5.506991
iteration 1200 / 1500: loss 5.668893
iteration 1300 / 1500: loss 5.570115
iteration 1400 / 1500: loss 5.794653
lr 1.000000e-07 reg 1.000000e+04 train accuracy: 0.385469 val accuracy: 0.391000
lr 1.000000e-07 reg 2.000000e+04 train accuracy: 0.372694 val accuracy: 0.394000
lr 1.000000e-07 reg 3.000000e+04 train accuracy: 0.364878 val accuracy: 0.366000
lr 2.000000e-07 reg 1.000000e+04 train accuracy: 0.373816 val accuracy: 0.387000
lr 2.000000e-07 reg 2.000000e+04 train accuracy: 0.364000 val accuracy: 0.370000
lr 2.000000e-07 reg 3.000000e+04 train accuracy: 0.363102 val accuracy: 0.376000
lr 4.000000e-07 reg 1.000000e+04 train accuracy: 0.369327 val accuracy: 0.363000
lr 4.000000e-07 reg 2.000000e+04 train accuracy: 0.345408 val accuracy: 0.367000
lr 4.000000e-07 reg 3.000000e+04 train accuracy: 0.326265 val accuracy: 0.345000
best validation accuracy achieved during cross-validation: 0.394000
(1e-07, 20000.0)
```

In [20]:

```
# Visualize the cross-validation results
import math
x_scatter = [math.log10(x[0]) for x in results]
y_scatter = [math.log10(x[1]) for x in results]

# plot training accuracy
marker_size = 100
colors = [results[x][0] for x in results]
plt.subplot(2, 1, 1)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 training accuracy')

# plot validation accuracy
colors = [results[x][1] for x in results] # default size of markers is 20
plt.subplot(2, 1, 2)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 validation accuracy')
plt.show()
```



In [21]:

```
# Evaluate the best svm on test set
y_test_pred = best_svm.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)
```

linear SVM on raw pixels final test set accuracy: 0.372000

In [22]:

```
# Visualize the learned weights for each class.
# Depending on your choice of learning rate and regularization strength, these may
# or may not be nice to look at.
w = best_svm.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)
w_min, w_max = np.min(w), np.max(w)
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```



## Inline question 2:

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look the way that they do.

**Your answer:** The visualized SVM weights look like blurry versions of each of their respective classes. They represent the average pixel values of each class (that maximize dot product with examples of this class).

# Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

In [1]:

```
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

In [2]:

```
def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
```

```

# subsample the data
mask = list(range(num_training, num_training + num_validation))
X_val = X_train[mask]
y_val = y_train[mask]
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Cleaning up variables to prevent loading data multiple times (which may cause
memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data(
)
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)

```

```
print('Test labels shape: ', y_test.shape)
```

```
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)
```

```
Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

## Softmax Classifier

Your code for this section will all be written inside **cs231n/classifiers/softmax.py**.

In [4]:

```
# First implement the naive softmax loss function with nested loops.
# Open the file cs231n/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from cs231n.classifiers.softmax import softmax_loss_naive, softmax_loss_vectorized
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)
#loss, grad = softmax_loss_vectorized(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))
```

```
loss: 2.337193
sanity check: 2.302585
```

## Inline Question 1:

Why do we expect our loss to be close to  $-\log(0.1)$ ? Explain briefly.\*\*

**Your answer:** There are 10 classes, and we initialized the weight matrix randomly s.t. each of the 10 classes should receive approximately the same score. Then  $\text{softmax}(x_i) = e^{\text{score}_{y_i}} / \sum e^{\text{score}_j} \approx 1/10$ ; softmax loss is average of  $-\log(\text{softmax}(x_i))$ .



In [5]:

```
# Complete the implementation of softmax_loss_naive and implement a (naive)  
# version of the gradient that uses nested loops.  
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)  
  
# As we did for the SVM, use numeric gradient checking as a debugging tool.  
# The numeric gradient should be close to the analytic gradient.  
from cs231n.gradient_check import grad_check_sparse  
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]  
grad_numerical = grad_check_sparse(f, W, grad, 10)  
  
# similar to SVM case, do another gradient check with regularization  
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)  
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]  
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: 0.567849 analytic: 0.567849, relative error: 1.022578e-07  
numerical: 1.100079 analytic: 1.100079, relative error: 2.135127e-08  
numerical: -2.462063 analytic: -2.462063, relative error: 2.582816e-  
09  
numerical: 2.206644 analytic: 2.206644, relative error: 1.724082e-08  
numerical: 3.106253 analytic: 3.106253, relative error: 1.736512e-08  
numerical: -1.265471 analytic: -1.265471, relative error: 1.625567e-  
08  
numerical: -0.510681 analytic: -0.510681, relative error: 6.360931e-  
08  
numerical: 1.432267 analytic: 1.432267, relative error: 5.457080e-08  
numerical: 0.851139 analytic: 0.851139, relative error: 8.969135e-10  
numerical: 1.133305 analytic: 1.133304, relative error: 4.096189e-08  
numerical: 1.635839 analytic: 1.635840, relative error: 2.319247e-08  
numerical: 0.500283 analytic: 0.500283, relative error: 1.672209e-08  
numerical: 1.967745 analytic: 1.967745, relative error: 3.375964e-08  
numerical: -1.071423 analytic: -1.071423, relative error: 3.551398e-  
08  
numerical: -1.091325 analytic: -1.091325, relative error: 5.064061e-  
08  
numerical: -2.049836 analytic: -2.049836, relative error: 3.272740e-  
08  
numerical: 0.887626 analytic: 0.887626, relative error: 5.251666e-09  
numerical: 2.146585 analytic: 2.146585, relative error: 2.256034e-08  
numerical: 2.850963 analytic: 2.850963, relative error: 9.449120e-09  
numerical: -0.237759 analytic: -0.237759, relative error: 8.487378e-  
08
```

In [6]:

```
# Now that we have a naive implementation of the softmax loss function and its gradient,  
# implement a vectorized version in softmax_loss_vectorized.  
# The two versions should compute the same results, but the vectorized version should be  
# much faster.  
tic = time.time()  
loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)  
toc = time.time()  
print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))  
  
from cs231n.classifiers.softmax import softmax_loss_vectorized  
tic = time.time()  
loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.000005)  
toc = time.time()  
print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))  
  
# As we did for the SVM, we use the Frobenius norm to compare the two versions of the gradient.  
grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')  
print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))  
print('Gradient difference: %f' % grad_difference)
```

```
naive loss: 2.337193e+00 computed in 0.071997s  
vectorized loss: 2.337193e+00 computed in 0.022640s  
Loss difference: 0.000000  
Gradient difference: 0.000000
```

In [15]:

```
# Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rates and regularization strengths; if you are careful you should be able to
# get a classification accuracy of over 0.35 on the validation set.
from cs231n.classifiers import Softmax
results = {}
best_val = -1
best_softmax = None
best_params = None
learning_rates = [5e-8, 1e-7, 5e-7]
regularization_strengths = [1e4, 2.5e4, 5e4]

for l in learning_rates:
    for r in regularization_strengths:
        soft = Softmax()
        curr_loss = soft.train(X_train, y_train, learning_rate=l, reg=r,
                               num_iters=1500, verbose=True)
        y_train_pred = soft.predict(X_train)
        y_train_acc = np.mean(y_train == y_train_pred)
        y_val_pred = soft.predict(X_val)
        y_val_acc = np.mean(y_val == y_val_pred)
        results[(l, r)] = (y_train_acc, y_val_acc)
        if y_val_acc > best_val:
            best_softmax = soft
            best_val = y_val_acc
            best_params = (l, r)

#####
#                               END OF YOUR CODE                               #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val
)
print(best_params)

iteration 0 / 1500: loss 314.378890
iteration 100 / 1500: loss 256.343739
iteration 200 / 1500: loss 210.005278
iteration 300 / 1500: loss 171.735479
iteration 400 / 1500: loss 140.718794
iteration 500 / 1500: loss 115.403647
iteration 600 / 1500: loss 94.755554
iteration 700 / 1500: loss 77.763083
iteration 800 / 1500: loss 63.903976
iteration 900 / 1500: loss 52.885643
```

iteration 1000 / 1500: loss 43.416895  
iteration 1100 / 1500: loss 35.993682  
iteration 1200 / 1500: loss 29.782689  
iteration 1300 / 1500: loss 24.691988  
iteration 1400 / 1500: loss 20.494863  
iteration 0 / 1500: loss 784.983024  
iteration 100 / 1500: loss 474.939836  
iteration 200 / 1500: loss 288.524443  
iteration 300 / 1500: loss 175.513011  
iteration 400 / 1500: loss 106.944714  
iteration 500 / 1500: loss 65.541919  
iteration 600 / 1500: loss 40.464141  
iteration 700 / 1500: loss 25.454078  
iteration 800 / 1500: loss 16.140550  
iteration 900 / 1500: loss 10.592509  
iteration 1000 / 1500: loss 7.199630  
iteration 1100 / 1500: loss 5.176507  
iteration 1200 / 1500: loss 3.956560  
iteration 1300 / 1500: loss 3.203671  
iteration 1400 / 1500: loss 2.765133  
iteration 0 / 1500: loss 1540.514962  
iteration 100 / 1500: loss 565.337485  
iteration 200 / 1500: loss 208.378826  
iteration 300 / 1500: loss 77.735027  
iteration 400 / 1500: loss 29.853285  
iteration 500 / 1500: loss 12.270621  
iteration 600 / 1500: loss 5.854505  
iteration 700 / 1500: loss 3.491227  
iteration 800 / 1500: loss 2.644319  
iteration 900 / 1500: loss 2.365326  
iteration 1000 / 1500: loss 2.205371  
iteration 1100 / 1500: loss 2.186416  
iteration 1200 / 1500: loss 2.182300  
iteration 1300 / 1500: loss 2.171551  
iteration 1400 / 1500: loss 2.079532  
iteration 0 / 1500: loss 310.737400  
iteration 100 / 1500: loss 206.861984  
iteration 200 / 1500: loss 138.510751  
iteration 300 / 1500: loss 93.337344  
iteration 400 / 1500: loss 63.009215  
iteration 500 / 1500: loss 42.864810  
iteration 600 / 1500: loss 29.195303  
iteration 700 / 1500: loss 20.230915  
iteration 800 / 1500: loss 14.230522  
iteration 900 / 1500: loss 10.132090  
iteration 1000 / 1500: loss 7.520559  
iteration 1100 / 1500: loss 5.596569  
iteration 1200 / 1500: loss 4.496602  
iteration 1300 / 1500: loss 3.632488  
iteration 1400 / 1500: loss 3.178314  
iteration 0 / 1500: loss 771.884038  
iteration 100 / 1500: loss 283.140275  
iteration 200 / 1500: loss 104.975657

iteration 300 / 1500: loss 39.794251  
iteration 400 / 1500: loss 15.861718  
iteration 500 / 1500: loss 7.124673  
iteration 600 / 1500: loss 3.922526  
iteration 700 / 1500: loss 2.728889  
iteration 800 / 1500: loss 2.305585  
iteration 900 / 1500: loss 2.191911  
iteration 1000 / 1500: loss 2.108062  
iteration 1100 / 1500: loss 2.040261  
iteration 1200 / 1500: loss 2.122895  
iteration 1300 / 1500: loss 2.086892  
iteration 1400 / 1500: loss 2.070165  
iteration 0 / 1500: loss 1569.807557  
iteration 100 / 1500: loss 211.542831  
iteration 200 / 1500: loss 30.125867  
iteration 300 / 1500: loss 5.896750  
iteration 400 / 1500: loss 2.632760  
iteration 500 / 1500: loss 2.200729  
iteration 600 / 1500: loss 2.157806  
iteration 700 / 1500: loss 2.166693  
iteration 800 / 1500: loss 2.115097  
iteration 900 / 1500: loss 2.131973  
iteration 1000 / 1500: loss 2.153319  
iteration 1100 / 1500: loss 2.122642  
iteration 1200 / 1500: loss 2.152296  
iteration 1300 / 1500: loss 2.166793  
iteration 1400 / 1500: loss 2.107457  
iteration 0 / 1500: loss 312.407420  
iteration 100 / 1500: loss 42.712501  
iteration 200 / 1500: loss 7.417511  
iteration 300 / 1500: loss 2.716093  
iteration 400 / 1500: loss 2.092821  
iteration 500 / 1500: loss 2.029881  
iteration 600 / 1500: loss 2.016174  
iteration 700 / 1500: loss 2.095447  
iteration 800 / 1500: loss 1.936520  
iteration 900 / 1500: loss 2.098716  
iteration 1000 / 1500: loss 2.050053  
iteration 1100 / 1500: loss 1.981984  
iteration 1200 / 1500: loss 2.037626  
iteration 1300 / 1500: loss 1.978531  
iteration 1400 / 1500: loss 1.960792  
iteration 0 / 1500: loss 783.468668  
iteration 100 / 1500: loss 6.962788  
iteration 200 / 1500: loss 2.155492  
iteration 300 / 1500: loss 2.089235  
iteration 400 / 1500: loss 2.045556  
iteration 500 / 1500: loss 2.068149  
iteration 600 / 1500: loss 2.066446  
iteration 700 / 1500: loss 2.076635  
iteration 800 / 1500: loss 2.057205  
iteration 900 / 1500: loss 2.138281  
iteration 1000 / 1500: loss 2.095712

```

iteration 1100 / 1500: loss 2.062105
iteration 1200 / 1500: loss 2.111504
iteration 1300 / 1500: loss 2.067248
iteration 1400 / 1500: loss 2.052028
iteration 0 / 1500: loss 1557.437666
iteration 100 / 1500: loss 2.220458
iteration 200 / 1500: loss 2.141399
iteration 300 / 1500: loss 2.189879
iteration 400 / 1500: loss 2.147661
iteration 500 / 1500: loss 2.144514
iteration 600 / 1500: loss 2.096290
iteration 700 / 1500: loss 2.167907
iteration 800 / 1500: loss 2.161713
iteration 900 / 1500: loss 2.168270
iteration 1000 / 1500: loss 2.180228
iteration 1100 / 1500: loss 2.179335
iteration 1200 / 1500: loss 2.163926
iteration 1300 / 1500: loss 2.086831
iteration 1400 / 1500: loss 2.185150
lr 5.000000e-08 reg 1.000000e+04 train accuracy: 0.297510 val accuracy: 0.303000
lr 5.000000e-08 reg 2.500000e+04 train accuracy: 0.326735 val accuracy: 0.331000
lr 5.000000e-08 reg 5.000000e+04 train accuracy: 0.311776 val accuracy: 0.327000
lr 1.000000e-07 reg 1.000000e+04 train accuracy: 0.353204 val accuracy: 0.366000
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.329041 val accuracy: 0.337000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.307714 val accuracy: 0.329000
lr 5.000000e-07 reg 1.000000e+04 train accuracy: 0.350714 val accuracy: 0.365000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.325653 val accuracy: 0.341000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.301816 val accuracy: 0.315000
best validation accuracy achieved during cross-validation: 0.366000
(1e-07, 10000.0)

```

In [16]:

```

# evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

```

```
softmax on raw pixels final test set accuracy: 0.354000
```

### Inline Question - True or False

It's possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your answer: True

Your explanation: With SVM loss, as long as the new image is predicted with the correct label over all other labels by greater than the margin, there will be an additional loss of 0. With softmax, as we predict the correct class with more confidence the additional loss approaches 0 ( $e^{\text{score\_yi}} \rightarrow \infty$ ,  $\log(e^{\text{score\_yi}} / \sum e^{\text{score\_j}} \rightarrow 1)$ ) but never truly reaches.

In [17]:


```
# Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)


classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship',
           'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```


plane




car




bird




cat




deer




dog




frog




horse



ship



truck





# Implementing a Neural Network

In this exercise we will develop a neural network with fully-connected layers to perform classification, and test it out on the CIFAR-10 dataset.

In [8]:

```
# A bit of setup

import numpy as np
import matplotlib.pyplot as plt

from cs231n.classifiers.neural_net import TwoLayerNet
from cs231n.classifiers.softmax import softmax_loss_vectorized

from __future__ import print_function

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

We will use the class `TwoLayerNet` in the file `cs231n/classifiers/neural_net.py` to represent instances of our network. The network parameters are stored in the instance variable `self.params` where keys are string parameter names and values are numpy arrays. Below, we initialize toy data and a toy model that we will use to develop your implementation.

In [9]:

```
# Create a small net and some toy data to check your implementations.
# Note that we set the random seed for repeatable experiments.

input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model():
    np.random.seed(0)
    return TwoLayerNet(input_size, hidden_size, num_classes, std=1e-1)

def init_toy_data():
    np.random.seed(1)
    X = 10 * np.random.randn(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

net = init_toy_model()
X, y = init_toy_data()
```

## Forward pass: compute scores

Open the file `cs231n/classifiers/neural_net.py` and look at the method `TwoLayerNet.loss`. This function is very similar to the loss functions you have written for the SVM and Softmax exercises: It takes the data and weights and computes the class scores, the loss, and the gradients on the parameters.

Implement the first part of the forward pass which uses the weights and biases to compute the scores for all inputs.

In [10]:

```
scores = net.loss(X)
print('Your scores:')
print(scores)
print()
print('correct scores:')
correct_scores = np.asarray([
    [-0.81233741, -1.27654624, -0.70335995],
    [-0.17129677, -1.18803311, -0.47310444],
    [-0.51590475, -1.01354314, -0.8504215 ],
    [-0.15419291, -0.48629638, -0.52901952],
    [-0.00618733, -0.12435261, -0.15226949]])
print(correct_scores)
print()

# The difference should be very small. We get < 1e-7
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))
```

Your scores:

```
[[-0.81233741 -1.27654624 -0.70335995]
 [-0.17129677 -1.18803311 -0.47310444]
 [-0.51590475 -1.01354314 -0.8504215 ]
 [-0.15419291 -0.48629638 -0.52901952]
 [-0.00618733 -0.12435261 -0.15226949]]
```

correct scores:

```
[[-0.81233741 -1.27654624 -0.70335995]
 [-0.17129677 -1.18803311 -0.47310444]
 [-0.51590475 -1.01354314 -0.8504215 ]
 [-0.15419291 -0.48629638 -0.52901952]
 [-0.00618733 -0.12435261 -0.15226949]]
```

Difference between your scores and correct scores:

```
3.6802720745909845e-08
```

## Forward pass: compute loss

In the same function, implement the second part that computes the data and regularization loss.

In [11]:

```
loss, _ = net.loss(X, y, reg=0.05)
correct_loss = 1.30378789133

# should be very small, we get < 1e-12
print('Difference between your loss and correct loss:')
print(np.sum(np.abs(loss - correct_loss)))
```

Difference between your loss and correct loss:  
1.7963408538435033e-13

## Backward pass

Implement the rest of the function. This will compute the gradient of the loss with respect to the variables  $w_1$ ,  $b_1$ ,  $w_2$ , and  $b_2$ . Now that you (hopefully!) have a correctly implemented forward pass, you can debug your backward pass using a numeric gradient check:

In [12]:

```
from cs231n.gradient_check import eval_numerical_gradient

# Use numeric gradient checking to check your implementation of the backward pass.
# If your implementation is correct, the difference between the numeric and
# analytic gradients should be less than 1e-8 for each of  $w_1$ ,  $w_2$ ,  $b_1$ , and  $b_2$ .

loss, grads = net.loss(X, y, reg=0.05)

# these should all be less than 1e-8 or so
for param_name in grads:
    f = lambda W: net.loss(X, y, reg=0.05)[0]
    param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
    print('%s max relative error: %e' % (param_name, rel_error(param_grad_num, grads[param_name])))
```

$w_2$  max relative error: 3.440708e-09  
 $b_2$  max relative error: 3.865112e-11  
 $w_1$  max relative error: 3.561318e-09  
 $b_1$  max relative error: 1.555471e-09

# Train the network

To train the network we will use stochastic gradient descent (SGD), similar to the SVM and Softmax classifiers. Look at the function `TwoLayerNet.train` and fill in the missing sections to implement the training procedure. This should be very similar to the training procedure you used for the SVM and Softmax classifiers. You will also have to implement `TwoLayerNet.predict`, as the training process periodically performs prediction to keep track of accuracy over time while the network trains.

Once you have implemented the method, run the code below to train a two-layer network on toy data. You should achieve a training loss less than 0.2.


In [7]:

```
net = init_toy_model()
stats = net.train(X, y, X, y,
                  learning_rate=1e-1, reg=5e-6,
                  num_iters=100, verbose=False)

print('Final training loss: ', stats['loss_history'][-1])

# plot the loss history
plt.plot(stats['loss_history'])
plt.xlabel('iteration')
plt.ylabel('training loss')
plt.title('Training Loss history')
plt.show()
```

Final training loss: 0.017149607938732093



## Load the data

Now that you have implemented a two-layer network that passes gradient checks and works on toy data, it's time to load up our favorite CIFAR-10 data so we can use it to train a classifier on a real dataset.

In [8]:

```
from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier. These are the same steps as
    we used for the SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
```

```

mask = list(range(num_training, num_training + num_validation))

X_val = X_train[mask]
y_val = y_train[mask]
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]

# Normalize the data: subtract the mean image
mean_image = np.mean(X_train, axis=0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image

# Reshape data to rows
X_train = X_train.reshape(num_training, -1)
X_val = X_val.reshape(num_validation, -1)
X_test = X_test.reshape(num_test, -1)

return X_train, y_train, X_val, y_val, X_test, y_test

# Cleaning up variables to prevent loading data multiple times (which may cause
memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

```

```

Train data shape:  (49000, 3072)
Train labels shape:  (49000,)
Validation data shape:  (1000, 3072)
Validation labels shape:  (1000,)
Test data shape:  (1000, 3072)
Test labels shape:  (1000,)

```

# Train a network

To train our network we will use SGD. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

In [9]:

```
input_size = 32 * 32 * 3
hidden_size = 50
num_classes = 10
net = TwoLayerNet(input_size, hidden_size, num_classes)

# Train the network
stats = net.train(X_train, y_train, X_val, y_val,
                  num_iters=1000, batch_size=200,
                  learning_rate=1e-4, learning_rate_decay=0.95,
                  reg=0.25, verbose=True)

# Predict on the validation set
val_acc = (net.predict(X_val) == y_val).mean()
print('Validation accuracy: ', val_acc)
```

```
iteration 0 / 1000: loss 2.302954
iteration 100 / 1000: loss 2.302550
iteration 200 / 1000: loss 2.297648
iteration 300 / 1000: loss 2.259602
iteration 400 / 1000: loss 2.204170
iteration 500 / 1000: loss 2.118565
iteration 600 / 1000: loss 2.051535
iteration 700 / 1000: loss 1.988466
iteration 800 / 1000: loss 2.006591
iteration 900 / 1000: loss 1.951473
Validation accuracy: 0.287
```

## Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.29 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.


Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.



In [10]:

```
# Plot the loss function and train / validation accuracies
plt.subplot(2, 1, 1)
plt.plot(stats['loss_history'])
plt.title('Loss history')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
plt.plot(stats['train_acc_history'], label='train')
plt.plot(stats['val_acc_history'], label='val')
plt.title('Classification accuracy history')
plt.xlabel('Epoch')
plt.ylabel('Classification accuracy')
plt.legend()
plt.show()
```




```
In [11]:
```

```
from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net.params['W1']
    W1 = W1.reshape(32, 32, 3, -1).transpose(3, 0, 1, 2)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(net)
```



# Tune your hyperparameters

**What's wrong?** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the learning rate decay, but you should be able to get good performance using the default value.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 48% on the validation set. Our best network gets over 52% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can, with a fully-connected Neural Network. Feel free to implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

In [23]:

```
#####  
#  
# TODO: Tune hyperparameters using the validation set. Store your best trained  
#  
# model in best_net.  
#  
#  
#  
# To help debug your network, it may help to use visualizations similar to the  
#  
# ones we used above; these visualizations will have significant qualitative  
#  
# differences from the ones we saw above for the poorly tuned network.  
#  
#  
#  
# Tweaking hyperparameters by hand can be fun, but you might find it useful to  
#  
# write code to sweep through possible combinations of hyperparameters  
#  
# automatically like we did on the previous exercises.  
#  
#####  
#  
# Your code  
#####  
#
```

```

#                                     END OF YOUR CODE

#####
#

results = {}
best_val = -1
best_net = None
best_param = None
learning_rates = [1e-3, 8e-4, 6e-4]
regularization_strengths = [0.4, 0.5, 0.6, 0.7, 0.8]

for l in learning_rates:
    for r in regularization_strengths:
        net = TwoLayerNet(input_size, hidden_size, num_classes)
        curr_loss = net.train(X_train, y_train, X_val, y_val, learning_rate=l, r
eg=r, num_iters=1500)
        y_train_pred = net.predict(X_train)
        y_train_acc = np.mean(y_train == y_train_pred)
        #print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
        y_val_pred = net.predict(X_val)
        y_val_acc = np.mean(y_val == y_val_pred)
        #print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))
        results[(l, r)] = (y_train_acc, y_val_acc)
        if y_val_acc > best_val:
            best_net = net
            best_val = y_val_acc
            best_params = (l, r)

#####
#                                     END OF YOUR CODE                                     #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val
)
print(best_params)

```

lr 6.000000e-04 reg 4.000000e-01 train accuracy: 0.483612 val accuracy: 0.476000  
lr 6.000000e-04 reg 5.000000e-01 train accuracy: 0.481449 val accuracy: 0.461000  
lr 6.000000e-04 reg 6.000000e-01 train accuracy: 0.474714 val accuracy: 0.455000  
lr 6.000000e-04 reg 7.000000e-01 train accuracy: 0.470959 val accuracy: 0.463000  
lr 6.000000e-04 reg 8.000000e-01 train accuracy: 0.471449 val accuracy: 0.460000  
lr 8.000000e-04 reg 4.000000e-01 train accuracy: 0.501959 val accuracy: 0.493000  
lr 8.000000e-04 reg 5.000000e-01 train accuracy: 0.494918 val accuracy: 0.476000  
lr 8.000000e-04 reg 6.000000e-01 train accuracy: 0.491082 val accuracy: 0.479000  
lr 8.000000e-04 reg 7.000000e-01 train accuracy: 0.488184 val accuracy: 0.465000  
lr 8.000000e-04 reg 8.000000e-01 train accuracy: 0.481857 val accuracy: 0.479000  
lr 1.000000e-03 reg 4.000000e-01 train accuracy: 0.500918 val accuracy: 0.479000  
lr 1.000000e-03 reg 5.000000e-01 train accuracy: 0.495286 val accuracy: 0.476000  
lr 1.000000e-03 reg 6.000000e-01 train accuracy: 0.492388 val accuracy: 0.468000  
lr 1.000000e-03 reg 7.000000e-01 train accuracy: 0.491714 val accuracy: 0.469000  
lr 1.000000e-03 reg 8.000000e-01 train accuracy: 0.493122 val accuracy: 0.468000  
best validation accuracy achieved during cross-validation: 0.493000  
(0.0008, 0.4)

In [24]:

```
# visualize the weights of the best network  
show_net_weights(best_net)
```



## Run on the test set

When you are done experimenting, you should evaluate your final trained network on the test set; you should get above 48%.

In [25]:

```
test_acc = (best_net.predict(X_test) == y_test).mean()  
print('Test accuracy: ', test_acc)
```

Test accuracy: 0.481

### Inline Question

Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

*Your answer:* 1, 3

*Your explanation:*

Testing accuracy lower than training accuracy means the model is overfitting. More data and regularization decreases the effect of specific datapoints / weights. More hidden units increases the complexity of the model which is more prone to overfit.

# Image features exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

We have seen that we can achieve reasonable performance on an image classification task by training a linear classifier on the pixels of the input image. In this exercise we will show that we can improve our classification performance by training linear classifiers not on raw pixels but on features that are computed from the raw pixels.

All of your work for this exercise will be done in this notebook.

In [26]:

```
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

## Load data

Similar to previous exercises, we will load CIFAR-10 data from disk.



In [27]:

```
from cs231n.features import color_histogram_hsv, hog_feature

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]

    return X_train, y_train, X_val, y_val, X_test, y_test

# Cleaning up variables to prevent loading data multiple times (which may cause
memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
```

Clear previously loaded data.

# Extract Features

For each image we will compute a Histogram of Oriented Gradients (HOG) as well as a color histogram using the hue channel in HSV color space. We form our final feature vector for each image by concatenating the HOG and color histogram feature vectors.

Roughly speaking, HOG should capture the texture of the image while ignoring color information, and the color histogram represents the color of the input image while ignoring texture. As a result, we expect that using both together ought to work better than using either alone. Verifying this assumption would be a good thing to try for your interests.

The `hog_feature` and `color_histogram_hsv` functions both operate on a single image and return a feature vector for that image. The `extract_features` function takes a set of images and a list of feature functions and evaluates each feature function on each image, storing the results in a matrix where each column is the concatenation of all feature vectors for a single image.

In [28]:

```
from cs231n.features import *

# 10, 11,
num_color_bins = 16 # Number of bins in the color histogram
feature_fns = [hog_feature, lambda img: color_histogram_hsv(img, nbin=num_color_bins)]
X_train_feats = extract_features(X_train, feature_fns, verbose=True)
X_val_feats = extract_features(X_val, feature_fns)
X_test_feats = extract_features(X_test, feature_fns)

# Preprocessing: Subtract the mean feature
mean_feat = np.mean(X_train_feats, axis=0, keepdims=True)
X_train_feats -= mean_feat
X_val_feats -= mean_feat
X_test_feats -= mean_feat

# Preprocessing: Divide by standard deviation. This ensures that each feature
# has roughly the same scale.
std_feat = np.std(X_train_feats, axis=0, keepdims=True)
X_train_feats /= std_feat
X_val_feats /= std_feat
X_test_feats /= std_feat

# Preprocessing: Add a bias dimension
X_train_feats = np.hstack([X_train_feats, np.ones((X_train_feats.shape[0], 1))])
X_val_feats = np.hstack([X_val_feats, np.ones((X_val_feats.shape[0], 1))])
X_test_feats = np.hstack([X_test_feats, np.ones((X_test_feats.shape[0], 1))])
```

[illegible]

## Train SVM on features

Using the multiclass SVM code developed earlier in the assignment, train SVMs on top of the features extracted above; this should achieve better results than training SVMs directly on top of raw pixels.

In [13]:

```
# Use the validation set to tune the learning rate and regularization strength

from cs231n.classifiers.linear_classifier import LinearSVM

learning_rates = [5e-8, 1e-7, 2e-7, 3e-7, 4e-7, 5e-7, 1e-6, 5e-6]
regularization_strengths = [1e3, 5e3, 1e4, 2e4, 3e4, 4e4, 5e4, 1e5]

results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation rate
.
best_params = None

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained classifier in best_svm. You might also want to play #
# with different numbers of bins in the color histogram. If you are careful #
# you should be able to get accuracy of near 0.44 on the validation set. #
#####
for l in learning_rates:
    for r in regularization_strengths:
        svm = LinearSVM()
        curr_loss = svm.train(X_train_feats, y_train, learning_rate=l, reg=r,
                               num_iters=1500, verbose=True)
        y_train_pred = svm.predict(X_train_feats)
        y_train_acc = np.mean(y_train == y_train_pred)
        y_val_pred = svm.predict(X_val_feats)
        y_val_acc = np.mean(y_val == y_val_pred)
        results[(l, r)] = (y_train_acc, y_val_acc)
        if y_val_acc > best_val:
            best_svm = svm
            best_val = y_val_acc
            best_params = (l, r)

#####
#                                     END OF YOUR CODE                                     #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val
)
print best_params
```

iteration 0 / 1500. loss 10.596329

```
iteration 0 / 1500: loss 10.390329
iteration 100 / 1500: loss 10.578521
iteration 200 / 1500: loss 10.546396
iteration 300 / 1500: loss 10.499155
iteration 400 / 1500: loss 10.469077
iteration 500 / 1500: loss 10.447657
iteration 600 / 1500: loss 10.414133
iteration 700 / 1500: loss 10.401073
iteration 800 / 1500: loss 10.364895
iteration 900 / 1500: loss 10.327907
iteration 1000 / 1500: loss 10.310513
iteration 1100 / 1500: loss 10.287885
iteration 1200 / 1500: loss 10.259661
iteration 1300 / 1500: loss 10.233150
iteration 1400 / 1500: loss 10.191958
iteration 0 / 1500: loss 16.635261
iteration 100 / 1500: loss 15.920372
iteration 200 / 1500: loss 15.261402
iteration 300 / 1500: loss 14.666318
iteration 400 / 1500: loss 14.126514
iteration 500 / 1500: loss 13.648507
iteration 600 / 1500: loss 13.197804
iteration 700 / 1500: loss 12.798200
iteration 800 / 1500: loss 12.440185
iteration 900 / 1500: loss 12.117257
iteration 1000 / 1500: loss 11.824610
iteration 1100 / 1500: loss 11.544930
iteration 1200 / 1500: loss 11.300652
iteration 1300 / 1500: loss 11.081145
iteration 1400 / 1500: loss 10.881673
iteration 0 / 1500: loss 24.597960
iteration 100 / 1500: loss 21.769524
iteration 200 / 1500: loss 19.453502
iteration 300 / 1500: loss 17.565047
iteration 400 / 1500: loss 16.009198
iteration 500 / 1500: loss 14.734748
iteration 600 / 1500: loss 13.695182
iteration 700 / 1500: loss 12.845738
iteration 800 / 1500: loss 12.145728
iteration 900 / 1500: loss 11.574274
iteration 1000 / 1500: loss 11.106844
iteration 1100 / 1500: loss 10.726786
iteration 1200 / 1500: loss 10.413845
iteration 1300 / 1500: loss 10.156120
iteration 1400 / 1500: loss 9.942285
iteration 0 / 1500: loss 41.215168
iteration 100 / 1500: loss 30.586929
iteration 200 / 1500: loss 23.459343
iteration 300 / 1500: loss 18.687359
iteration 400 / 1500: loss 15.492110
iteration 500 / 1500: loss 13.353427
iteration 600 / 1500: loss 11.912268
iteration 700 / 1500: loss 10.955144
iteration 800 / 1500: loss 10.307624
```

iteration 900 / 1500: loss 9.876250  
iteration 1000 / 1500: loss 9.587906  
iteration 1100 / 1500: loss 9.392915  
iteration 1200 / 1500: loss 9.264058  
iteration 1300 / 1500: loss 9.175502  
iteration 1400 / 1500: loss 9.116656  
iteration 0 / 1500: loss 56.992314  
iteration 100 / 1500: loss 35.318036  
iteration 200 / 1500: loss 23.428484  
iteration 300 / 1500: loss 16.912829  
iteration 400 / 1500: loss 13.338948  
iteration 500 / 1500: loss 11.377607  
iteration 600 / 1500: loss 10.304441  
iteration 700 / 1500: loss 9.713867  
iteration 800 / 1500: loss 9.392074  
iteration 900 / 1500: loss 9.214761  
iteration 1000 / 1500: loss 9.117002  
iteration 1100 / 1500: loss 9.064144  
iteration 1200 / 1500: loss 9.034676  
iteration 1300 / 1500: loss 9.018974  
iteration 1400 / 1500: loss 9.009795  
iteration 0 / 1500: loss 71.806144  
iteration 100 / 1500: loss 37.177618  
iteration 200 / 1500: loss 21.640609  
iteration 300 / 1500: loss 14.671909  
iteration 400 / 1500: loss 11.544970  
iteration 500 / 1500: loss 10.142297  
iteration 600 / 1500: loss 9.512133  
iteration 700 / 1500: loss 9.229340  
iteration 800 / 1500: loss 9.102448  
iteration 900 / 1500: loss 9.045769  
iteration 1000 / 1500: loss 9.020374  
iteration 1100 / 1500: loss 9.009014  
iteration 1200 / 1500: loss 9.003781  
iteration 1300 / 1500: loss 9.001471  
iteration 1400 / 1500: loss 9.000446  
iteration 0 / 1500: loss 88.089470  
iteration 100 / 1500: loss 38.027785  
iteration 200 / 1500: loss 19.648601  
iteration 300 / 1500: loss 12.908922  
iteration 400 / 1500: loss 10.437089  
iteration 500 / 1500: loss 9.525816  
iteration 600 / 1500: loss 9.192766  
iteration 700 / 1500: loss 9.070484  
iteration 800 / 1500: loss 9.025501  
iteration 900 / 1500: loss 9.009187  
iteration 1000 / 1500: loss 9.003053  
iteration 1100 / 1500: loss 9.000979  
iteration 1200 / 1500: loss 9.000162  
iteration 1300 / 1500: loss 8.999800  
iteration 1400 / 1500: loss 8.999701  
iteration 0 / 1500: loss 171.061714  
iteration 100 / 1500: loss 30.716381

iteration 200 / 1500: loss 11.909215  
iteration 300 / 1500: loss 9.389579  
iteration 400 / 1500: loss 9.052185  
iteration 500 / 1500: loss 9.006856  
iteration 600 / 1500: loss 9.000763  
iteration 700 / 1500: loss 8.999907  
iteration 800 / 1500: loss 8.999812  
iteration 900 / 1500: loss 8.999824  
iteration 1000 / 1500: loss 8.999816  
iteration 1100 / 1500: loss 8.999858  
iteration 1200 / 1500: loss 8.999816  
iteration 1300 / 1500: loss 8.999841  
iteration 1400 / 1500: loss 8.999859  
iteration 0 / 1500: loss 10.649769  
iteration 100 / 1500: loss 10.601363  
iteration 200 / 1500: loss 10.510027  
iteration 300 / 1500: loss 10.454438  
iteration 400 / 1500: loss 10.401835  
iteration 500 / 1500: loss 10.344445  
iteration 600 / 1500: loss 10.300307  
iteration 700 / 1500: loss 10.260208  
iteration 800 / 1500: loss 10.193068  
iteration 900 / 1500: loss 10.161515  
iteration 1000 / 1500: loss 10.106179  
iteration 1100 / 1500: loss 10.050118  
iteration 1200 / 1500: loss 10.007931  
iteration 1300 / 1500: loss 9.982270  
iteration 1400 / 1500: loss 9.938138  
iteration 0 / 1500: loss 17.387181  
iteration 100 / 1500: loss 15.866039  
iteration 200 / 1500: loss 14.617982  
iteration 300 / 1500: loss 13.600652  
iteration 400 / 1500: loss 12.771689  
iteration 500 / 1500: loss 12.079819  
iteration 600 / 1500: loss 11.524240  
iteration 700 / 1500: loss 11.064045  
iteration 800 / 1500: loss 10.686178  
iteration 900 / 1500: loss 10.379494  
iteration 1000 / 1500: loss 10.134963  
iteration 1100 / 1500: loss 9.930588  
iteration 1200 / 1500: loss 9.757500  
iteration 1300 / 1500: loss 9.618199  
iteration 1400 / 1500: loss 9.504882  
iteration 0 / 1500: loss 24.795060  
iteration 100 / 1500: loss 19.577447  
iteration 200 / 1500: loss 16.091228  
iteration 300 / 1500: loss 13.756382  
iteration 400 / 1500: loss 12.181977  
iteration 500 / 1500: loss 11.134599  
iteration 600 / 1500: loss 10.432101  
iteration 700 / 1500: loss 9.959688  
iteration 800 / 1500: loss 9.640130  
iteration 900 / 1500: loss 9.429855



iteration 1000 / 1500: loss 9.287308  
iteration 1100 / 1500: loss 9.190578  
iteration 1200 / 1500: loss 9.125790  
iteration 1300 / 1500: loss 9.084592  
iteration 1400 / 1500: loss 9.056440  
iteration 0 / 1500: loss 41.984678  
iteration 100 / 1500: loss 23.792966  
iteration 200 / 1500: loss 15.632379  
iteration 300 / 1500: loss 11.978102  
iteration 400 / 1500: loss 10.336052  
iteration 500 / 1500: loss 9.596862  
iteration 600 / 1500: loss 9.268282  
iteration 700 / 1500: loss 9.120058  
iteration 800 / 1500: loss 9.053599  
iteration 900 / 1500: loss 9.022866  
iteration 1000 / 1500: loss 9.009820  
iteration 1100 / 1500: loss 9.003661  
iteration 1200 / 1500: loss 9.001228  
iteration 1300 / 1500: loss 9.000169  
iteration 1400 / 1500: loss 8.999427  
iteration 0 / 1500: loss 54.622317  
iteration 100 / 1500: loss 22.694929  
iteration 200 / 1500: loss 13.106683  
iteration 300 / 1500: loss 10.235222  
iteration 400 / 1500: loss 9.369835  
iteration 500 / 1500: loss 9.110756  
iteration 600 / 1500: loss 9.032846  
iteration 700 / 1500: loss 9.009534  
iteration 800 / 1500: loss 9.002205  
iteration 900 / 1500: loss 9.000426  
iteration 1000 / 1500: loss 8.999723  
iteration 1100 / 1500: loss 8.999557  
iteration 1200 / 1500: loss 8.999541  
iteration 1300 / 1500: loss 8.999424  
iteration 1400 / 1500: loss 8.999423  
iteration 0 / 1500: loss 73.772330  
iteration 100 / 1500: loss 21.988922  
iteration 200 / 1500: loss 11.608119  
iteration 300 / 1500: loss 9.521271  
iteration 400 / 1500: loss 9.104405  
iteration 500 / 1500: loss 9.020349  
iteration 600 / 1500: loss 9.003797  
iteration 700 / 1500: loss 9.000310  
iteration 800 / 1500: loss 8.999705  
iteration 900 / 1500: loss 8.999640  
iteration 1000 / 1500: loss 8.999554  
iteration 1100 / 1500: loss 8.999570  
iteration 1200 / 1500: loss 8.999551  
iteration 1300 / 1500: loss 8.999696  
iteration 1400 / 1500: loss 8.999587  
iteration 0 / 1500: loss 87.630900  
iteration 100 / 1500: loss 19.528651  
iteration 200 / 1500: loss 10.410475

iteration 300 / 1500: loss 9.188930  
iteration 400 / 1500: loss 9.024899  
iteration 500 / 1500: loss 9.003082  
iteration 600 / 1500: loss 9.000108  
iteration 700 / 1500: loss 8.999695  
iteration 800 / 1500: loss 8.999671  
iteration 900 / 1500: loss 8.999654  
iteration 1000 / 1500: loss 8.999692  
iteration 1100 / 1500: loss 8.999623  
iteration 1200 / 1500: loss 8.999664  
iteration 1300 / 1500: loss 8.999582  
iteration 1400 / 1500: loss 8.999663  
iteration 0 / 1500: loss 169.096718  
iteration 100 / 1500: loss 11.814801  
iteration 200 / 1500: loss 9.049346  
iteration 300 / 1500: loss 9.000670  
iteration 400 / 1500: loss 8.999870  
iteration 500 / 1500: loss 8.999839  
iteration 600 / 1500: loss 8.999842  
iteration 700 / 1500: loss 8.999862  
iteration 800 / 1500: loss 8.999791  
iteration 900 / 1500: loss 8.999797  
iteration 1000 / 1500: loss 8.999793  
iteration 1100 / 1500: loss 8.999786  
iteration 1200 / 1500: loss 8.999813  
iteration 1300 / 1500: loss 8.999809  
iteration 1400 / 1500: loss 8.999826  
iteration 0 / 1500: loss 10.604011  
iteration 100 / 1500: loss 10.472676  
iteration 200 / 1500: loss 10.364896  
iteration 300 / 1500: loss 10.261710  
iteration 400 / 1500: loss 10.167433  
iteration 500 / 1500: loss 10.074861  
iteration 600 / 1500: loss 9.979947  
iteration 700 / 1500: loss 9.914838  
iteration 800 / 1500: loss 9.828529  
iteration 900 / 1500: loss 9.781575  
iteration 1000 / 1500: loss 9.710182  
iteration 1100 / 1500: loss 9.670637  
iteration 1200 / 1500: loss 9.611616  
iteration 1300 / 1500: loss 9.563630  
iteration 1400 / 1500: loss 9.509502  
iteration 0 / 1500: loss 16.966789  
iteration 100 / 1500: loss 14.350894  
iteration 200 / 1500: loss 12.579941  
iteration 300 / 1500: loss 11.394810  
iteration 400 / 1500: loss 10.613195  
iteration 500 / 1500: loss 10.076459  
iteration 600 / 1500: loss 9.720416  
iteration 700 / 1500: loss 9.480665  
iteration 800 / 1500: loss 9.319702  
iteration 900 / 1500: loss 9.213474  
iteration 1000 / 1500: loss 9.142741

iteration 1100 / 1500: loss 9.094373  
iteration 1200 / 1500: loss 9.061986  
iteration 1300 / 1500: loss 9.040417  
iteration 1400 / 1500: loss 9.025286  
iteration 0 / 1500: loss 25.081834  
iteration 100 / 1500: loss 16.212045  
iteration 200 / 1500: loss 12.230897  
iteration 300 / 1500: loss 10.446150  
iteration 400 / 1500: loss 9.647403  
iteration 500 / 1500: loss 9.290825  
iteration 600 / 1500: loss 9.128026  
iteration 700 / 1500: loss 9.057532  
iteration 800 / 1500: loss 9.024519  
iteration 900 / 1500: loss 9.010335  
iteration 1000 / 1500: loss 9.003811  
iteration 1100 / 1500: loss 9.000578  
iteration 1200 / 1500: loss 8.999315  
iteration 1300 / 1500: loss 8.998705  
iteration 1400 / 1500: loss 8.998336  
iteration 0 / 1500: loss 41.144295  
iteration 100 / 1500: loss 15.447825  
iteration 200 / 1500: loss 10.292492  
iteration 300 / 1500: loss 9.259019  
iteration 400 / 1500: loss 9.051028  
iteration 500 / 1500: loss 9.009489  
iteration 600 / 1500: loss 9.001051  
iteration 700 / 1500: loss 8.999504  
iteration 800 / 1500: loss 8.999237  
iteration 900 / 1500: loss 8.999150  
iteration 1000 / 1500: loss 8.999052  
iteration 1100 / 1500: loss 8.999022  
iteration 1200 / 1500: loss 8.999202  
iteration 1300 / 1500: loss 8.999341  
iteration 1400 / 1500: loss 8.999178  
iteration 0 / 1500: loss 58.600333  
iteration 100 / 1500: loss 13.433920  
iteration 200 / 1500: loss 9.395946  
iteration 300 / 1500: loss 9.034582  
iteration 400 / 1500: loss 9.002482  
iteration 500 / 1500: loss 8.999674  
iteration 600 / 1500: loss 8.999379  
iteration 700 / 1500: loss 8.999499  
iteration 800 / 1500: loss 8.999400  
iteration 900 / 1500: loss 8.999474  
iteration 1000 / 1500: loss 8.999220  
iteration 1100 / 1500: loss 8.999395  
iteration 1200 / 1500: loss 8.999348  
iteration 1300 / 1500: loss 8.999499  
iteration 1400 / 1500: loss 8.999431  
iteration 0 / 1500: loss 75.681033  
iteration 100 / 1500: loss 11.651563  
iteration 200 / 1500: loss 9.104614  
iteration 300 / 1500: loss 9.003786

iteration 400 / 1500: loss 8.999635  
iteration 500 / 1500: loss 8.999590  
iteration 600 / 1500: loss 8.999578  
iteration 700 / 1500: loss 8.999654  
iteration 800 / 1500: loss 8.999546  
iteration 900 / 1500: loss 8.999539  
iteration 1000 / 1500: loss 8.999626  
iteration 1100 / 1500: loss 8.999588  
iteration 1200 / 1500: loss 8.999535  
iteration 1300 / 1500: loss 8.999426  
iteration 1400 / 1500: loss 8.999521  
iteration 0 / 1500: loss 90.584749  
iteration 100 / 1500: loss 10.432575  
iteration 200 / 1500: loss 9.025119  
iteration 300 / 1500: loss 9.000137  
iteration 400 / 1500: loss 8.999540  
iteration 500 / 1500: loss 8.999557  
iteration 600 / 1500: loss 8.999643  
iteration 700 / 1500: loss 8.999586  
iteration 800 / 1500: loss 8.999671  
iteration 900 / 1500: loss 8.999660  
iteration 1000 / 1500: loss 8.999664  
iteration 1100 / 1500: loss 8.999622  
iteration 1200 / 1500: loss 8.999616  
iteration 1300 / 1500: loss 8.999588  
iteration 1400 / 1500: loss 8.999640  
iteration 0 / 1500: loss 173.694308  
iteration 100 / 1500: loss 9.046457  
iteration 200 / 1500: loss 8.999853  
iteration 300 / 1500: loss 8.999820  
iteration 400 / 1500: loss 8.999819  
iteration 500 / 1500: loss 8.999787  
iteration 600 / 1500: loss 8.999841  
iteration 700 / 1500: loss 8.999827  
iteration 800 / 1500: loss 8.999826  
iteration 900 / 1500: loss 8.999828  
iteration 1000 / 1500: loss 8.999815  
iteration 1100 / 1500: loss 8.999838  
iteration 1200 / 1500: loss 8.999837  
iteration 1300 / 1500: loss 8.999824  
iteration 1400 / 1500: loss 8.999834  
iteration 0 / 1500: loss 10.597976  
iteration 100 / 1500: loss 10.422787  
iteration 200 / 1500: loss 10.261017  
iteration 300 / 1500: loss 10.116479  
iteration 400 / 1500: loss 9.975429  
iteration 500 / 1500: loss 9.873461  
iteration 600 / 1500: loss 9.773045  
iteration 700 / 1500: loss 9.676795  
iteration 800 / 1500: loss 9.597652  
iteration 900 / 1500: loss 9.521525  
iteration 1000 / 1500: loss 9.463032  
iteration 1100 / 1500: loss 9.420537

iteration 1200 / 1500: loss 9.366744  
iteration 1300 / 1500: loss 9.323150  
iteration 1400 / 1500: loss 9.282847  
iteration 0 / 1500: loss 17.233463  
iteration 100 / 1500: loss 13.505898  
iteration 200 / 1500: loss 11.479238  
iteration 300 / 1500: loss 10.357619  
iteration 400 / 1500: loss 9.741784  
iteration 500 / 1500: loss 9.404683  
iteration 600 / 1500: loss 9.222597  
iteration 700 / 1500: loss 9.120459  
iteration 800 / 1500: loss 9.062282  
iteration 900 / 1500: loss 9.032998  
iteration 1000 / 1500: loss 9.016069  
iteration 1100 / 1500: loss 9.007595  
iteration 1200 / 1500: loss 9.002323  
iteration 1300 / 1500: loss 8.999683  
iteration 1400 / 1500: loss 8.998030  
iteration 0 / 1500: loss 25.399315  
iteration 100 / 1500: loss 13.916744  
iteration 200 / 1500: loss 10.472596  
iteration 300 / 1500: loss 9.441882  
iteration 400 / 1500: loss 9.129961  
iteration 500 / 1500: loss 9.038050  
iteration 600 / 1500: loss 9.010227  
iteration 700 / 1500: loss 9.001496  
iteration 800 / 1500: loss 8.999173  
iteration 900 / 1500: loss 8.998699  
iteration 1000 / 1500: loss 8.997935  
iteration 1100 / 1500: loss 8.998308  
iteration 1200 / 1500: loss 8.998244  
iteration 1300 / 1500: loss 8.998004  
iteration 1400 / 1500: loss 8.998457  
iteration 0 / 1500: loss 41.763473  
iteration 100 / 1500: loss 11.928966  
iteration 200 / 1500: loss 9.260535  
iteration 300 / 1500: loss 9.022207  
iteration 400 / 1500: loss 9.001227  
iteration 500 / 1500: loss 8.999098  
iteration 600 / 1500: loss 8.999064  
iteration 700 / 1500: loss 8.999171  
iteration 800 / 1500: loss 8.999094  
iteration 900 / 1500: loss 8.999110  
iteration 1000 / 1500: loss 8.999224  
iteration 1100 / 1500: loss 8.999188  
iteration 1200 / 1500: loss 8.999206  
iteration 1300 / 1500: loss 8.999165  
iteration 1400 / 1500: loss 8.998876  
iteration 0 / 1500: loss 57.657941  
iteration 100 / 1500: loss 10.287605  
iteration 200 / 1500: loss 9.033523  
iteration 300 / 1500: loss 9.000453  
iteration 400 / 1500: loss 8.999401

iteration 500 / 1500: loss 8.999362  
iteration 600 / 1500: loss 8.999563  
iteration 700 / 1500: loss 8.999465  
iteration 800 / 1500: loss 8.999545  
iteration 900 / 1500: loss 8.999398  
iteration 1000 / 1500: loss 8.999354  
iteration 1100 / 1500: loss 8.999484  
iteration 1200 / 1500: loss 8.999383  
iteration 1300 / 1500: loss 8.999403  
iteration 1400 / 1500: loss 8.999410  
iteration 0 / 1500: loss 75.892270  
iteration 100 / 1500: loss 9.520302  
iteration 200 / 1500: loss 9.003729  
iteration 300 / 1500: loss 8.999623  
iteration 400 / 1500: loss 8.999502  
iteration 500 / 1500: loss 8.999656  
iteration 600 / 1500: loss 8.999631  
iteration 700 / 1500: loss 8.999546  
iteration 800 / 1500: loss 8.999476  
iteration 900 / 1500: loss 8.999507  
iteration 1000 / 1500: loss 8.999578  
iteration 1100 / 1500: loss 8.999559  
iteration 1200 / 1500: loss 8.999521  
iteration 1300 / 1500: loss 8.999553  
iteration 1400 / 1500: loss 8.999527  
iteration 0 / 1500: loss 89.518852  
iteration 100 / 1500: loss 9.181799  
iteration 200 / 1500: loss 9.000124  
iteration 300 / 1500: loss 8.999649  
iteration 400 / 1500: loss 8.999701  
iteration 500 / 1500: loss 8.999619  
iteration 600 / 1500: loss 8.999628  
iteration 700 / 1500: loss 8.999640  
iteration 800 / 1500: loss 8.999594  
iteration 900 / 1500: loss 8.999611  
iteration 1000 / 1500: loss 8.999734  
iteration 1100 / 1500: loss 8.999569  
iteration 1200 / 1500: loss 8.999629  
iteration 1300 / 1500: loss 8.999596  
iteration 1400 / 1500: loss 8.999583  
iteration 0 / 1500: loss 161.811135  
iteration 100 / 1500: loss 9.000467  
iteration 200 / 1500: loss 8.999852  
iteration 300 / 1500: loss 8.999852  
iteration 400 / 1500: loss 8.999797  
iteration 500 / 1500: loss 8.999783  
iteration 600 / 1500: loss 8.999849  
iteration 700 / 1500: loss 8.999827  
iteration 800 / 1500: loss 8.999826  
iteration 900 / 1500: loss 8.999832  
iteration 1000 / 1500: loss 8.999842  
iteration 1100 / 1500: loss 8.999863  
iteration 1200 / 1500: loss 8.999851

iteration 1300 / 1500: loss 8.999799  
iteration 1400 / 1500: loss 8.999831  
iteration 0 / 1500: loss 10.729373  
iteration 100 / 1500: loss 10.465829  
iteration 200 / 1500: loss 10.264722  
iteration 300 / 1500: loss 10.064632  
iteration 400 / 1500: loss 9.909230  
iteration 500 / 1500: loss 9.758948  
iteration 600 / 1500: loss 9.651990  
iteration 700 / 1500: loss 9.559395  
iteration 800 / 1500: loss 9.472006  
iteration 900 / 1500: loss 9.403993  
iteration 1000 / 1500: loss 9.333529  
iteration 1100 / 1500: loss 9.281952  
iteration 1200 / 1500: loss 9.235019  
iteration 1300 / 1500: loss 9.203810  
iteration 1400 / 1500: loss 9.165482  
iteration 0 / 1500: loss 17.327822  
iteration 100 / 1500: loss 12.731024  
iteration 200 / 1500: loss 10.664237  
iteration 300 / 1500: loss 9.747534  
iteration 400 / 1500: loss 9.337824  
iteration 500 / 1500: loss 9.147404  
iteration 600 / 1500: loss 9.063699  
iteration 700 / 1500: loss 9.027401  
iteration 800 / 1500: loss 9.010290  
iteration 900 / 1500: loss 9.002423  
iteration 1000 / 1500: loss 8.998465  
iteration 1100 / 1500: loss 8.996756  
iteration 1200 / 1500: loss 8.996777  
iteration 1300 / 1500: loss 8.996360  
iteration 1400 / 1500: loss 8.996326  
iteration 0 / 1500: loss 24.396710  
iteration 100 / 1500: loss 12.090010  
iteration 200 / 1500: loss 9.618799  
iteration 300 / 1500: loss 9.122673  
iteration 400 / 1500: loss 9.023045  
iteration 500 / 1500: loss 9.003485  
iteration 600 / 1500: loss 8.999619  
iteration 700 / 1500: loss 8.998247  
iteration 800 / 1500: loss 8.998359  
iteration 900 / 1500: loss 8.997749  
iteration 1000 / 1500: loss 8.998128  
iteration 1100 / 1500: loss 8.998368  
iteration 1200 / 1500: loss 8.998521  
iteration 1300 / 1500: loss 8.998490  
iteration 1400 / 1500: loss 8.998483  
iteration 0 / 1500: loss 42.884748  
iteration 100 / 1500: loss 10.346457  
iteration 200 / 1500: loss 9.052748  
iteration 300 / 1500: loss 9.001272  
iteration 400 / 1500: loss 8.999069  
iteration 500 / 1500: loss 8.998997

iteration 600 / 1500: loss 8.999062  
iteration 700 / 1500: loss 8.999167  
iteration 800 / 1500: loss 8.999007  
iteration 900 / 1500: loss 8.999224  
iteration 1000 / 1500: loss 8.999115  
iteration 1100 / 1500: loss 8.999282  
iteration 1200 / 1500: loss 8.999243  
iteration 1300 / 1500: loss 8.999141  
iteration 1400 / 1500: loss 8.999122  
iteration 0 / 1500: loss 57.042659  
iteration 100 / 1500: loss 9.372103  
iteration 200 / 1500: loss 9.002303  
iteration 300 / 1500: loss 8.999517  
iteration 400 / 1500: loss 8.999503  
iteration 500 / 1500: loss 8.999315  
iteration 600 / 1500: loss 8.999387  
iteration 700 / 1500: loss 8.999357  
iteration 800 / 1500: loss 8.999399  
iteration 900 / 1500: loss 8.999391  
iteration 1000 / 1500: loss 8.999494  
iteration 1100 / 1500: loss 8.999353  
iteration 1200 / 1500: loss 8.999349  
iteration 1300 / 1500: loss 8.999480  
iteration 1400 / 1500: loss 8.999298  
iteration 0 / 1500: loss 72.398296  
iteration 100 / 1500: loss 9.094401  
iteration 200 / 1500: loss 8.999673  
iteration 300 / 1500: loss 8.999496  
iteration 400 / 1500: loss 8.999612  
iteration 500 / 1500: loss 8.999604  
iteration 600 / 1500: loss 8.999550  
iteration 700 / 1500: loss 8.999600  
iteration 800 / 1500: loss 8.999535  
iteration 900 / 1500: loss 8.999512  
iteration 1000 / 1500: loss 8.999676  
iteration 1100 / 1500: loss 8.999496  
iteration 1200 / 1500: loss 8.999633  
iteration 1300 / 1500: loss 8.999523  
iteration 1400 / 1500: loss 8.999590  
iteration 0 / 1500: loss 90.342942  
iteration 100 / 1500: loss 9.022753  
iteration 200 / 1500: loss 8.999724  
iteration 300 / 1500: loss 8.999726  
iteration 400 / 1500: loss 8.999644  
iteration 500 / 1500: loss 8.999647  
iteration 600 / 1500: loss 8.999721  
iteration 700 / 1500: loss 8.999716  
iteration 800 / 1500: loss 8.999638  
iteration 900 / 1500: loss 8.999684  
iteration 1000 / 1500: loss 8.999610  
iteration 1100 / 1500: loss 8.999708  
iteration 1200 / 1500: loss 8.999627  
iteration 1300 / 1500: loss 8.999707



iteration 1400 / 1500: loss 8.999655  
iteration 0 / 1500: loss 173.556081  
iteration 100 / 1500: loss 8.999862  
iteration 200 / 1500: loss 8.999848  
iteration 300 / 1500: loss 8.999832  
iteration 400 / 1500: loss 8.999807  
iteration 500 / 1500: loss 8.999828  
iteration 600 / 1500: loss 8.999817  
iteration 700 / 1500: loss 8.999855  
iteration 800 / 1500: loss 8.999844  
iteration 900 / 1500: loss 8.999846  
iteration 1000 / 1500: loss 8.999830  
iteration 1100 / 1500: loss 8.999879  
iteration 1200 / 1500: loss 8.999808  
iteration 1300 / 1500: loss 8.999812  
iteration 1400 / 1500: loss 8.999814  
iteration 0 / 1500: loss 10.570393  
iteration 100 / 1500: loss 10.264311  
iteration 200 / 1500: loss 10.047732  
iteration 300 / 1500: loss 9.855832  
iteration 400 / 1500: loss 9.684709  
iteration 500 / 1500: loss 9.565026  
iteration 600 / 1500: loss 9.451793  
iteration 700 / 1500: loss 9.367290  
iteration 800 / 1500: loss 9.297359  
iteration 900 / 1500: loss 9.247584  
iteration 1000 / 1500: loss 9.196811  
iteration 1100 / 1500: loss 9.159311  
iteration 1200 / 1500: loss 9.122963  
iteration 1300 / 1500: loss 9.105164  
iteration 1400 / 1500: loss 9.079747  
iteration 0 / 1500: loss 16.981030  
iteration 100 / 1500: loss 11.920799  
iteration 200 / 1500: loss 10.069927  
iteration 300 / 1500: loss 9.392087  
iteration 400 / 1500: loss 9.143255  
iteration 500 / 1500: loss 9.048390  
iteration 600 / 1500: loss 9.015748  
iteration 700 / 1500: loss 9.004352  
iteration 800 / 1500: loss 8.999161  
iteration 900 / 1500: loss 8.997283  
iteration 1000 / 1500: loss 8.997040  
iteration 1100 / 1500: loss 8.996720  
iteration 1200 / 1500: loss 8.996523  
iteration 1300 / 1500: loss 8.997189  
iteration 1400 / 1500: loss 8.996433  
iteration 0 / 1500: loss 25.307496  
iteration 100 / 1500: loss 11.188444  
iteration 200 / 1500: loss 9.293045  
iteration 300 / 1500: loss 9.037573  
iteration 400 / 1500: loss 9.003320  
iteration 500 / 1500: loss 8.999025  
iteration 600 / 1500: loss 8.998391

iteration 700 / 1500: loss 8.998195  
iteration 800 / 1500: loss 8.998018  
iteration 900 / 1500: loss 8.998440  
iteration 1000 / 1500: loss 8.998252  
iteration 1100 / 1500: loss 8.998526  
iteration 1200 / 1500: loss 8.998471  
iteration 1300 / 1500: loss 8.997971  
iteration 1400 / 1500: loss 8.998179  
iteration 0 / 1500: loss 41.368160  
iteration 100 / 1500: loss 9.569448  
iteration 200 / 1500: loss 9.008812  
iteration 300 / 1500: loss 8.999039  
iteration 400 / 1500: loss 8.999124  
iteration 500 / 1500: loss 8.999072  
iteration 600 / 1500: loss 8.999272  
iteration 700 / 1500: loss 8.998984  
iteration 800 / 1500: loss 8.999159  
iteration 900 / 1500: loss 8.998962  
iteration 1000 / 1500: loss 8.999010  
iteration 1100 / 1500: loss 8.999078  
iteration 1200 / 1500: loss 8.999084  
iteration 1300 / 1500: loss 8.999181  
iteration 1400 / 1500: loss 8.999083  
iteration 0 / 1500: loss 54.815418  
iteration 100 / 1500: loss 9.102831  
iteration 200 / 1500: loss 8.999669  
iteration 300 / 1500: loss 8.999447  
iteration 400 / 1500: loss 8.999405  
iteration 500 / 1500: loss 8.999452  
iteration 600 / 1500: loss 8.999530  
iteration 700 / 1500: loss 8.999408  
iteration 800 / 1500: loss 8.999338  
iteration 900 / 1500: loss 8.999496  
iteration 1000 / 1500: loss 8.999509  
iteration 1100 / 1500: loss 8.999203  
iteration 1200 / 1500: loss 8.999437  
iteration 1300 / 1500: loss 8.999345  
iteration 1400 / 1500: loss 8.999330  
iteration 0 / 1500: loss 75.362154  
iteration 100 / 1500: loss 9.018300  
iteration 200 / 1500: loss 8.999580  
iteration 300 / 1500: loss 8.999650  
iteration 400 / 1500: loss 8.999458  
iteration 500 / 1500: loss 8.999642  
iteration 600 / 1500: loss 8.999487  
iteration 700 / 1500: loss 8.999515  
iteration 800 / 1500: loss 8.999649  
iteration 900 / 1500: loss 8.999542  
iteration 1000 / 1500: loss 8.999501  
iteration 1100 / 1500: loss 8.999602  
iteration 1200 / 1500: loss 8.999582  
iteration 1300 / 1500: loss 8.999585  
iteration 1400 / 1500: loss 8.999560

iteration 0 / 1500: loss 86.802750  
iteration 100 / 1500: loss 9.002371  
iteration 200 / 1500: loss 8.999649  
iteration 300 / 1500: loss 8.999634  
iteration 400 / 1500: loss 8.999653  
iteration 500 / 1500: loss 8.999657  
iteration 600 / 1500: loss 8.999642  
iteration 700 / 1500: loss 8.999632  
iteration 800 / 1500: loss 8.999622  
iteration 900 / 1500: loss 8.999620  
iteration 1000 / 1500: loss 8.999635  
iteration 1100 / 1500: loss 8.999676  
iteration 1200 / 1500: loss 8.999669  
iteration 1300 / 1500: loss 8.999626  
iteration 1400 / 1500: loss 8.999655  
iteration 0 / 1500: loss 161.531972  
iteration 100 / 1500: loss 8.999815  
iteration 200 / 1500: loss 8.999822  
iteration 300 / 1500: loss 8.999835  
iteration 400 / 1500: loss 8.999802  
iteration 500 / 1500: loss 8.999814  
iteration 600 / 1500: loss 8.999784  
iteration 700 / 1500: loss 8.999858  
iteration 800 / 1500: loss 8.999833  
iteration 900 / 1500: loss 8.999802  
iteration 1000 / 1500: loss 8.999793  
iteration 1100 / 1500: loss 8.999823  
iteration 1200 / 1500: loss 8.999889  
iteration 1300 / 1500: loss 8.999873  
iteration 1400 / 1500: loss 8.999832  
iteration 0 / 1500: loss 10.587828  
iteration 100 / 1500: loss 10.052891  
iteration 200 / 1500: loss 9.702093  
iteration 300 / 1500: loss 9.458461  
iteration 400 / 1500: loss 9.306641  
iteration 500 / 1500: loss 9.196640  
iteration 600 / 1500: loss 9.128164  
iteration 700 / 1500: loss 9.076964  
iteration 800 / 1500: loss 9.044735  
iteration 900 / 1500: loss 9.026516  
iteration 1000 / 1500: loss 9.010755  
iteration 1100 / 1500: loss 9.001037  
iteration 1200 / 1500: loss 8.994874  
iteration 1300 / 1500: loss 8.989765  
iteration 1400 / 1500: loss 8.989583  
iteration 0 / 1500: loss 16.568270  
iteration 100 / 1500: loss 10.011221  
iteration 200 / 1500: loss 9.133680  
iteration 300 / 1500: loss 9.013916  
iteration 400 / 1500: loss 8.998662  
iteration 500 / 1500: loss 8.996388  
iteration 600 / 1500: loss 8.996028  
iteration 700 / 1500: loss 8.996395

iteration 800 / 1500: loss 8.996658  
iteration 900 / 1500: loss 8.995615  
iteration 1000 / 1500: loss 8.996330  
iteration 1100 / 1500: loss 8.996738  
iteration 1200 / 1500: loss 8.996399  
iteration 1300 / 1500: loss 8.996673  
iteration 1400 / 1500: loss 8.996325  
iteration 0 / 1500: loss 24.854267  
iteration 100 / 1500: loss 9.278244  
iteration 200 / 1500: loss 9.002978  
iteration 300 / 1500: loss 8.998335  
iteration 400 / 1500: loss 8.998720  
iteration 500 / 1500: loss 8.998517  
iteration 600 / 1500: loss 8.997834  
iteration 700 / 1500: loss 8.998193  
iteration 800 / 1500: loss 8.997918  
iteration 900 / 1500: loss 8.998207  
iteration 1000 / 1500: loss 8.998369  
iteration 1100 / 1500: loss 8.997978  
iteration 1200 / 1500: loss 8.998069  
iteration 1300 / 1500: loss 8.998181  
iteration 1400 / 1500: loss 8.998478  
iteration 0 / 1500: loss 41.654656  
iteration 100 / 1500: loss 9.008239  
iteration 200 / 1500: loss 8.998992  
iteration 300 / 1500: loss 8.998919  
iteration 400 / 1500: loss 8.999109  
iteration 500 / 1500: loss 8.999208  
iteration 600 / 1500: loss 8.999118  
iteration 700 / 1500: loss 8.998965  
iteration 800 / 1500: loss 8.999188  
iteration 900 / 1500: loss 8.998938  
iteration 1000 / 1500: loss 8.999082  
iteration 1100 / 1500: loss 8.999238  
iteration 1200 / 1500: loss 8.999082  
iteration 1300 / 1500: loss 8.999304  
iteration 1400 / 1500: loss 8.999151  
iteration 0 / 1500: loss 58.108084  
iteration 100 / 1500: loss 8.999538  
iteration 200 / 1500: loss 8.999393  
iteration 300 / 1500: loss 8.999497  
iteration 400 / 1500: loss 8.999308  
iteration 500 / 1500: loss 8.999375  
iteration 600 / 1500: loss 8.999478  
iteration 700 / 1500: loss 8.999545  
iteration 800 / 1500: loss 8.999515  
iteration 900 / 1500: loss 8.999319  
iteration 1000 / 1500: loss 8.999495  
iteration 1100 / 1500: loss 8.999382  
iteration 1200 / 1500: loss 8.999423  
iteration 1300 / 1500: loss 8.999396  
iteration 1400 / 1500: loss 8.999458  
iteration 0 / 1500: loss 75.523520

iteration 100 / 1500: loss 8.999494  
iteration 200 / 1500: loss 8.999553  
iteration 300 / 1500: loss 8.999547  
iteration 400 / 1500: loss 8.999507  
iteration 500 / 1500: loss 8.999614  
iteration 600 / 1500: loss 8.999510  
iteration 700 / 1500: loss 8.999666  
iteration 800 / 1500: loss 8.999523  
iteration 900 / 1500: loss 8.999602  
iteration 1000 / 1500: loss 8.999673  
iteration 1100 / 1500: loss 8.999549  
iteration 1200 / 1500: loss 8.999615  
iteration 1300 / 1500: loss 8.999557  
iteration 1400 / 1500: loss 8.999568  
iteration 0 / 1500: loss 85.872054  
iteration 100 / 1500: loss 8.999647  
iteration 200 / 1500: loss 8.999685  
iteration 300 / 1500: loss 8.999686  
iteration 400 / 1500: loss 8.999697  
iteration 500 / 1500: loss 8.999716  
iteration 600 / 1500: loss 8.999745  
iteration 700 / 1500: loss 8.999635  
iteration 800 / 1500: loss 8.999690  
iteration 900 / 1500: loss 8.999684  
iteration 1000 / 1500: loss 8.999669  
iteration 1100 / 1500: loss 8.999610  
iteration 1200 / 1500: loss 8.999605  
iteration 1300 / 1500: loss 8.999643  
iteration 1400 / 1500: loss 8.999704  
iteration 0 / 1500: loss 173.709683  
iteration 100 / 1500: loss 8.999843  
iteration 200 / 1500: loss 8.999824  
iteration 300 / 1500: loss 8.999836  
iteration 400 / 1500: loss 8.999850  
iteration 500 / 1500: loss 8.999813  
iteration 600 / 1500: loss 8.999864  
iteration 700 / 1500: loss 8.999886  
iteration 800 / 1500: loss 8.999812  
iteration 900 / 1500: loss 8.999853  
iteration 1000 / 1500: loss 8.999815  
iteration 1100 / 1500: loss 8.999854  
iteration 1200 / 1500: loss 8.999830  
iteration 1300 / 1500: loss 8.999827  
iteration 1400 / 1500: loss 8.999812  
iteration 0 / 1500: loss 10.644058  
iteration 100 / 1500: loss 9.207834  
iteration 200 / 1500: loss 9.011109  
iteration 300 / 1500: loss 8.988933  
iteration 400 / 1500: loss 8.984772  
iteration 500 / 1500: loss 8.983945  
iteration 600 / 1500: loss 8.983209  
iteration 700 / 1500: loss 8.980983  
iteration 800 / 1500: loss 8.981477

iteration 900 / 1500: loss 8.981896  
iteration 1000 / 1500: loss 8.979910  
iteration 1100 / 1500: loss 8.983869  
iteration 1200 / 1500: loss 8.979598  
iteration 1300 / 1500: loss 8.980929  
iteration 1400 / 1500: loss 8.983668  
iteration 0 / 1500: loss 17.379961  
iteration 100 / 1500: loss 8.996374  
iteration 200 / 1500: loss 8.995983  
iteration 300 / 1500: loss 8.995851  
iteration 400 / 1500: loss 8.997152  
iteration 500 / 1500: loss 8.996199  
iteration 600 / 1500: loss 8.996209  
iteration 700 / 1500: loss 8.995864  
iteration 800 / 1500: loss 8.995865  
iteration 900 / 1500: loss 8.996624  
iteration 1000 / 1500: loss 8.996295  
iteration 1100 / 1500: loss 8.995846  
iteration 1200 / 1500: loss 8.996987  
iteration 1300 / 1500: loss 8.996483  
iteration 1400 / 1500: loss 8.996271  
iteration 0 / 1500: loss 25.867173  
iteration 100 / 1500: loss 8.998509  
iteration 200 / 1500: loss 8.998708  
iteration 300 / 1500: loss 8.998210  
iteration 400 / 1500: loss 8.998283  
iteration 500 / 1500: loss 8.998032  
iteration 600 / 1500: loss 8.998144  
iteration 700 / 1500: loss 8.998215  
iteration 800 / 1500: loss 8.998542  
iteration 900 / 1500: loss 8.998573  
iteration 1000 / 1500: loss 8.997999  
iteration 1100 / 1500: loss 8.998693  
iteration 1200 / 1500: loss 8.998112  
iteration 1300 / 1500: loss 8.998575  
iteration 1400 / 1500: loss 8.997898  
iteration 0 / 1500: loss 42.014623  
iteration 100 / 1500: loss 8.999297  
iteration 200 / 1500: loss 8.998963  
iteration 300 / 1500: loss 8.999306  
iteration 400 / 1500: loss 8.999105  
iteration 500 / 1500: loss 8.999178  
iteration 600 / 1500: loss 8.999129  
iteration 700 / 1500: loss 8.999370  
iteration 800 / 1500: loss 8.999198  
iteration 900 / 1500: loss 8.999138  
iteration 1000 / 1500: loss 8.999369  
iteration 1100 / 1500: loss 8.999068  
iteration 1200 / 1500: loss 8.999063  
iteration 1300 / 1500: loss 8.999171  
iteration 1400 / 1500: loss 8.999220  
iteration 0 / 1500: loss 57.341200  
iteration 100 / 1500: loss 8.999549

iteration 200 / 1500: loss 8.999563  
iteration 300 / 1500: loss 8.999709  
iteration 400 / 1500: loss 8.999482  
iteration 500 / 1500: loss 8.999421  
iteration 600 / 1500: loss 8.999578  
iteration 700 / 1500: loss 8.999588  
iteration 800 / 1500: loss 8.999414  
iteration 900 / 1500: loss 8.999363  
iteration 1000 / 1500: loss 8.999493  
iteration 1100 / 1500: loss 8.999473  
iteration 1200 / 1500: loss 8.999486  
iteration 1300 / 1500: loss 8.999490  
iteration 1400 / 1500: loss 8.999658  
iteration 0 / 1500: loss 74.177490  
iteration 100 / 1500: loss 8.999791  
iteration 200 / 1500: loss 8.999668  
iteration 300 / 1500: loss 8.999729  
iteration 400 / 1500: loss 8.999738  
iteration 500 / 1500: loss 8.999627  
iteration 600 / 1500: loss 8.999656  
iteration 700 / 1500: loss 8.999694  
iteration 800 / 1500: loss 8.999696  
iteration 900 / 1500: loss 8.999695  
iteration 1000 / 1500: loss 8.999639  
iteration 1100 / 1500: loss 8.999769  
iteration 1200 / 1500: loss 8.999578  
iteration 1300 / 1500: loss 8.999685  
iteration 1400 / 1500: loss 8.999714  
iteration 0 / 1500: loss 93.935979  
iteration 100 / 1500: loss 8.999766  
iteration 200 / 1500: loss 8.999712  
iteration 300 / 1500: loss 8.999835  
iteration 400 / 1500: loss 8.999750  
iteration 500 / 1500: loss 8.999812  
iteration 600 / 1500: loss 8.999726  
iteration 700 / 1500: loss 8.999796  
iteration 800 / 1500: loss 8.999877  
iteration 900 / 1500: loss 8.999825  
iteration 1000 / 1500: loss 8.999736  
iteration 1100 / 1500: loss 8.999741  
iteration 1200 / 1500: loss 8.999696  
iteration 1300 / 1500: loss 8.999707  
iteration 1400 / 1500: loss 8.999687  
iteration 0 / 1500: loss 172.675160  
iteration 100 / 1500: loss 9.000009  
iteration 200 / 1500: loss 8.999969  
iteration 300 / 1500: loss 9.000000  
iteration 400 / 1500: loss 8.999967  
iteration 500 / 1500: loss 8.999976  
iteration 600 / 1500: loss 8.999971  
iteration 700 / 1500: loss 9.000044  
iteration 800 / 1500: loss 8.999956  
iteration 900 / 1500: loss 8.999927

iteration 1000 / 1500: loss 9.000006  
iteration 1100 / 1500: loss 8.999998  
iteration 1200 / 1500: loss 9.000052  
iteration 1300 / 1500: loss 9.000013  
iteration 1400 / 1500: loss 9.000010  
lr 5.000000e-08 reg 1.000000e+03 train accuracy: 0.098429 val accuracy: 0.098000  
lr 5.000000e-08 reg 5.000000e+03 train accuracy: 0.113735 val accuracy: 0.117000  
lr 5.000000e-08 reg 1.000000e+04 train accuracy: 0.132571 val accuracy: 0.123000  
lr 5.000000e-08 reg 2.000000e+04 train accuracy: 0.177735 val accuracy: 0.192000  
lr 5.000000e-08 reg 3.000000e+04 train accuracy: 0.246490 val accuracy: 0.239000  
lr 5.000000e-08 reg 4.000000e+04 train accuracy: 0.382469 val accuracy: 0.400000  
lr 5.000000e-08 reg 5.000000e+04 train accuracy: 0.416327 val accuracy: 0.408000  
lr 5.000000e-08 reg 1.000000e+05 train accuracy: 0.417878 val accuracy: 0.422000  
lr 1.000000e-07 reg 1.000000e+03 train accuracy: 0.123122 val accuracy: 0.134000  
lr 1.000000e-07 reg 5.000000e+03 train accuracy: 0.151082 val accuracy: 0.147000  
lr 1.000000e-07 reg 1.000000e+04 train accuracy: 0.216490 val accuracy: 0.228000  
lr 1.000000e-07 reg 2.000000e+04 train accuracy: 0.405327 val accuracy: 0.409000  
lr 1.000000e-07 reg 3.000000e+04 train accuracy: 0.417612 val accuracy: 0.419000  
lr 1.000000e-07 reg 4.000000e+04 train accuracy: 0.422959 val accuracy: 0.420000  
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.417020 val accuracy: 0.426000  
lr 1.000000e-07 reg 1.000000e+05 train accuracy: 0.423184 val accuracy: 0.435000  
lr 2.000000e-07 reg 1.000000e+03 train accuracy: 0.146694 val accuracy: 0.144000  
lr 2.000000e-07 reg 5.000000e+03 train accuracy: 0.291020 val accuracy: 0.305000  
lr 2.000000e-07 reg 1.000000e+04 train accuracy: 0.415959 val accuracy: 0.406000  
lr 2.000000e-07 reg 2.000000e+04 train accuracy: 0.417551 val accuracy: 0.428000  
lr 2.000000e-07 reg 3.000000e+04 train accuracy: 0.419061 val accuracy: 0.438000  
lr 2.000000e-07 reg 4.000000e+04 train accuracy: 0.416653 val accuracy: 0.412000  
lr 2.000000e-07 reg 5.000000e+04 train accuracy: 0.419898 val accuracy: 0.420000  
lr 2.000000e-07 reg 1.000000e+05 train accuracy: 0.412796 val accuracy: 0.404000



lr 3.000000e-07 reg 1.000000e+03 train accuracy: 0.205796 val accuracy: 0.216000  
lr 3.000000e-07 reg 5.000000e+03 train accuracy: 0.397898 val accuracy: 0.393000  
lr 3.000000e-07 reg 1.000000e+04 train accuracy: 0.418143 val accuracy: 0.426000  
lr 3.000000e-07 reg 2.000000e+04 train accuracy: 0.419245 val accuracy: 0.421000  
lr 3.000000e-07 reg 3.000000e+04 train accuracy: 0.420061 val accuracy: 0.429000  
lr 3.000000e-07 reg 4.000000e+04 train accuracy: 0.414653 val accuracy: 0.420000  
lr 3.000000e-07 reg 5.000000e+04 train accuracy: 0.412490 val accuracy: 0.413000  
lr 3.000000e-07 reg 1.000000e+05 train accuracy: 0.414673 val accuracy: 0.422000  
lr 4.000000e-07 reg 1.000000e+03 train accuracy: 0.246531 val accuracy: 0.249000  
lr 4.000000e-07 reg 5.000000e+03 train accuracy: 0.417490 val accuracy: 0.422000  
lr 4.000000e-07 reg 1.000000e+04 train accuracy: 0.418898 val accuracy: 0.421000  
lr 4.000000e-07 reg 2.000000e+04 train accuracy: 0.417347 val accuracy: 0.421000  
lr 4.000000e-07 reg 3.000000e+04 train accuracy: 0.416531 val accuracy: 0.410000  
lr 4.000000e-07 reg 4.000000e+04 train accuracy: 0.414816 val accuracy: 0.412000  
lr 4.000000e-07 reg 5.000000e+04 train accuracy: 0.408816 val accuracy: 0.404000  
lr 4.000000e-07 reg 1.000000e+05 train accuracy: 0.420000 val accuracy: 0.415000  
lr 5.000000e-07 reg 1.000000e+03 train accuracy: 0.261265 val accuracy: 0.303000  
lr 5.000000e-07 reg 5.000000e+03 train accuracy: 0.416816 val accuracy: 0.425000  
lr 5.000000e-07 reg 1.000000e+04 train accuracy: 0.420408 val accuracy: 0.432000  
lr 5.000000e-07 reg 2.000000e+04 train accuracy: 0.412306 val accuracy: 0.414000  
lr 5.000000e-07 reg 3.000000e+04 train accuracy: 0.416204 val accuracy: 0.407000  
lr 5.000000e-07 reg 4.000000e+04 train accuracy: 0.417796 val accuracy: 0.417000  
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.414102 val accuracy: 0.418000  
lr 5.000000e-07 reg 1.000000e+05 train accuracy: 0.407510 val accuracy: 0.405000  
lr 1.000000e-06 reg 1.000000e+03 train accuracy: 0.408980 val accuracy: 0.401000  
lr 1.000000e-06 reg 5.000000e+03 train accuracy: 0.416551 val accuracy: 0.415000  
lr 1.000000e-06 reg 1.000000e+04 train accuracy: 0.417694 val accuracy:

```
cy: 0.412000
lr 1.000000e-06 reg 2.000000e+04 train accuracy: 0.415306 val accuracy: 0.412000
lr 1.000000e-06 reg 3.000000e+04 train accuracy: 0.419918 val accuracy: 0.420000
lr 1.000000e-06 reg 4.000000e+04 train accuracy: 0.407918 val accuracy: 0.404000
lr 1.000000e-06 reg 5.000000e+04 train accuracy: 0.415796 val accuracy: 0.411000
lr 1.000000e-06 reg 1.000000e+05 train accuracy: 0.405612 val accuracy: 0.410000
lr 5.000000e-06 reg 1.000000e+03 train accuracy: 0.420694 val accuracy: 0.421000
lr 5.000000e-06 reg 5.000000e+03 train accuracy: 0.411082 val accuracy: 0.416000
lr 5.000000e-06 reg 1.000000e+04 train accuracy: 0.412694 val accuracy: 0.416000
lr 5.000000e-06 reg 2.000000e+04 train accuracy: 0.411633 val accuracy: 0.413000
lr 5.000000e-06 reg 3.000000e+04 train accuracy: 0.398939 val accuracy: 0.408000
lr 5.000000e-06 reg 4.000000e+04 train accuracy: 0.364837 val accuracy: 0.356000
lr 5.000000e-06 reg 5.000000e+04 train accuracy: 0.369918 val accuracy: 0.374000
lr 5.000000e-06 reg 1.000000e+05 train accuracy: 0.325653 val accuracy: 0.325000
best validation accuracy achieved during cross-validation: 0.438000
```

In [14]:

```
# Evaluate your trained SVM on the test set
y_test_pred = best_svm.predict(X_test_feats)
test_accuracy = np.mean(y_test == y_test_pred)
print(test_accuracy)
```

0.426

In [15]:

```
# An important way to gain intuition about how an algorithm works is to
# visualize the mistakes that it makes. In this visualization, we show examples
# of images that are misclassified by our current system. The first column
# shows images that our system labeled as "plane" but whose true label is
# something other than "plane".

examples_per_class = 8
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for cls, cls_name in enumerate(classes):
    idxs = np.where((y_test != cls) & (y_test_pred == cls))[0]
    idxs = np.random.choice(idxs, examples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt.subplot(examples_per_class, len(classes), i * len(classes) + cls + 1)

        plt.imshow(X_test[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls_name)
plt.show()
```



## Inline question 1:

Describe the misclassification results that you see. Do they make sense?

Some seem to make sense - for example, peripheral view of bird that can arguably look like a plane taking off or a dog on a grassy field standing on 4 legs mistaken for a horse. But many other misclassifications are hard to interpret.

# Neural Network on image features

Earlier in this assignment we saw that training a two-layer neural network on raw pixels achieved better classification performance than linear classifiers on raw pixels. In this notebook we have seen that linear classifiers on image features outperform linear classifiers on raw pixels.

For completeness, we should also try training a neural network on image features. This approach should outperform all previous approaches: you should easily be able to achieve over 55% classification accuracy on the test set; our best model achieves about 60% classification accuracy.

In [29]:

```
# Preprocessing: Remove the bias dimension  
# Make sure to run this cell only ONCE  
print(X_train_feats.shape)  
X_train_feats = X_train_feats[:, :-1]  
X_val_feats = X_val_feats[:, :-1]  
X_test_feats = X_test_feats[:, :-1]  
  
print(X_train_feats.shape)
```

```
(49000, 161)
```

```
(49000, 160)
```

In [33]:

```
from cs231n.classifiers.neural_net import TwoLayerNet

input_dim = X_train_feats.shape[1]
hidden_dim = 500
num_classes = 10

#####
# TODO: Train a two-layer neural network on image features. You may want to #
# cross-validate various parameters as in previous sections. Store your best #
# model in the best_net variable. #
#####
results = {}
best_val = -1
best_net = None
best_params = None
learning_rates = [1e-1, 5e-1, 1e0]
regularization_strengths = [1e-3, 5e-3, 1e-2]

for l in learning_rates:
    for r in regularization_strengths:
        print(l)
        net = TwoLayerNet(input_dim, hidden_dim, num_classes)
        curr_loss = net.train(X_train_feats, y_train, X_val_feats, y_val, learning_rate=l, reg=r, num_iters=1500)
        y_train_pred = net.predict(X_train_feats)
        y_train_acc = np.mean(y_train == y_train_pred)
        y_val_pred = net.predict(X_val_feats)
        y_val_acc = np.mean(y_val == y_val_pred)
        results[(l, r)] = (y_train_acc, y_val_acc)
        if y_val_acc > best_val:
            best_net = net
            best_val = y_val_acc
            best_params = (l, r)

#####
#                                     END OF YOUR CODE                                     #
#####
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)
print(best_params)
```

```
0.1
0.1
0.1
0.5
0.5
0.5
1.0
1.0
1.0
lr 1.000000e-01 reg 1.000000e-03 train accuracy: 0.540918 val accuracy: 0.525000
lr 1.000000e-01 reg 5.000000e-03 train accuracy: 0.528388 val accuracy: 0.514000
lr 1.000000e-01 reg 1.000000e-02 train accuracy: 0.521020 val accuracy: 0.520000
lr 5.000000e-01 reg 1.000000e-03 train accuracy: 0.660694 val accuracy: 0.597000
lr 5.000000e-01 reg 5.000000e-03 train accuracy: 0.565143 val accuracy: 0.559000
lr 5.000000e-01 reg 1.000000e-02 train accuracy: 0.514714 val accuracy: 0.500000
lr 1.000000e+00 reg 1.000000e-03 train accuracy: 0.671020 val accuracy: 0.560000
lr 1.000000e+00 reg 5.000000e-03 train accuracy: 0.555163 val accuracy: 0.519000
lr 1.000000e+00 reg 1.000000e-02 train accuracy: 0.509469 val accuracy: 0.486000
best validation accuracy achieved during cross-validation: 0.597000
(0.5, 0.001)
```

In [34]:

```
# Run your best neural net classifier on the test set. You should be able
# to get more than 55% accuracy.
```

```
test_acc = (best_net.predict(X_test_feats) == y_test).mean()
print(test_acc)
```

```
0.577
```