

# Image features exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

We have seen that we can achieve reasonable performance on an image classification task by training a linear classifier on the pixels of the input image. In this exercise we will show that we can improve our classification performance by training linear classifiers not on raw pixels but on features that are computed from the raw pixels.

All of your work for this exercise will be done in this notebook.

In [26]:

```
import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

from __future__ import print_function

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

## Load data

Similar to previous exercises, we will load CIFAR-10 data from disk.

In [27]:

```
from cs231n.features import color_histogram_hsv, hog_feature

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]

    return X_train, y_train, X_val, y_val, X_test, y_test

# Cleaning up variables to prevent loading data multiple times (which may cause
memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
```

Clear previously loaded data.

# Extract Features

For each image we will compute a Histogram of Oriented Gradients (HOG) as well as a color histogram using the hue channel in HSV color space. We form our final feature vector for each image by concatenating the HOG and color histogram feature vectors.

Roughly speaking, HOG should capture the texture of the image while ignoring color information, and the color histogram represents the color of the input image while ignoring texture. As a result, we expect that using both together ought to work better than using either alone. Verifying this assumption would be a good thing to try for your interests.

The `hog_feature` and `color_histogram_hsv` functions both operate on a single image and return a feature vector for that image. The `extract_features` function takes a set of images and a list of feature functions and evaluates each feature function on each image, storing the results in a matrix where each column is the concatenation of all feature vectors for a single image.

In [28]:

```
from cs231n.features import *

# 10, 11,
num_color_bins = 16 # Number of bins in the color histogram
feature_fns = [hog_feature, lambda img: color_histogram_hsv(img, nbin=num_color_
bins)]
X_train_feats = extract_features(X_train, feature_fns, verbose=True)
X_val_feats = extract_features(X_val, feature_fns)
X_test_feats = extract_features(X_test, feature_fns)

# Preprocessing: Subtract the mean feature
mean_feat = np.mean(X_train_feats, axis=0, keepdims=True)
X_train_feats -= mean_feat
X_val_feats -= mean_feat
X_test_feats -= mean_feat

# Preprocessing: Divide by standard deviation. This ensures that each feature
# has roughly the same scale.
std_feat = np.std(X_train_feats, axis=0, keepdims=True)
X_train_feats /= std_feat
X_val_feats /= std_feat
X_test_feats /= std_feat

# Preprocessing: Add a bias dimension
X_train_feats = np.hstack([X_train_feats, np.ones((X_train_feats.shape[0], 1))])
X_val_feats = np.hstack([X_val_feats, np.ones((X_val_feats.shape[0], 1))])
X_test_feats = np.hstack([X_test_feats, np.ones((X_test_feats.shape[0], 1))])
```

[illegible]

## Train SVM on features

Using the multiclass SVM code developed earlier in the assignment, train SVMs on top of the features extracted above; this should achieve better results than training SVMs directly on top of raw pixels.

In [13]:

```
# Use the validation set to tune the learning rate and regularization strength

from cs231n.classifiers.linear_classifier import LinearSVM

learning_rates = [5e-8, 1e-7, 2e-7, 3e-7, 4e-7, 5e-7, 1e-6, 5e-6]
regularization_strengths = [1e3, 5e3, 1e4, 2e4, 3e4, 4e4, 5e4, 1e5]

results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation rate
.
best_params = None

#####
# TODO: #
# Use the validation set to set the learning rate and regularization strength. #
# This should be identical to the validation that you did for the SVM; save #
# the best trained classifier in best_svm. You might also want to play #
# with different numbers of bins in the color histogram. If you are careful #
# you should be able to get accuracy of near 0.44 on the validation set. #
#####
for l in learning_rates:
    for r in regularization_strengths:
        svm = LinearSVM()
        curr_loss = svm.train(X_train_feats, y_train, learning_rate=l, reg=r,
                               num_iters=1500, verbose=True)
        y_train_pred = svm.predict(X_train_feats)
        y_train_acc = np.mean(y_train == y_train_pred)
        y_val_pred = svm.predict(X_val_feats)
        y_val_acc = np.mean(y_val == y_val_pred)
        results[(l, r)] = (y_train_acc, y_val_acc)
        if y_val_acc > best_val:
            best_svm = svm
            best_val = y_val_acc
            best_params = (l, r)

#####
#                               END OF YOUR CODE                               #
#####

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val
)
print best_params

iteration 0 / 1500: loss 10.596329
```

iteration 0 / 1500: loss 10.590529  
iteration 100 / 1500: loss 10.578521  
iteration 200 / 1500: loss 10.546396  
iteration 300 / 1500: loss 10.499155  
iteration 400 / 1500: loss 10.469077  
iteration 500 / 1500: loss 10.447657  
iteration 600 / 1500: loss 10.414133  
iteration 700 / 1500: loss 10.401073  
iteration 800 / 1500: loss 10.364895  
iteration 900 / 1500: loss 10.327907  
iteration 1000 / 1500: loss 10.310513  
iteration 1100 / 1500: loss 10.287885  
iteration 1200 / 1500: loss 10.259661  
iteration 1300 / 1500: loss 10.233150  
iteration 1400 / 1500: loss 10.191958  
iteration 0 / 1500: loss 16.635261  
iteration 100 / 1500: loss 15.920372  
iteration 200 / 1500: loss 15.261402  
iteration 300 / 1500: loss 14.666318  
iteration 400 / 1500: loss 14.126514  
iteration 500 / 1500: loss 13.648507  
iteration 600 / 1500: loss 13.197804  
iteration 700 / 1500: loss 12.798200  
iteration 800 / 1500: loss 12.440185  
iteration 900 / 1500: loss 12.117257  
iteration 1000 / 1500: loss 11.824610  
iteration 1100 / 1500: loss 11.544930  
iteration 1200 / 1500: loss 11.300652  
iteration 1300 / 1500: loss 11.081145  
iteration 1400 / 1500: loss 10.881673  
iteration 0 / 1500: loss 24.597960  
iteration 100 / 1500: loss 21.769524  
iteration 200 / 1500: loss 19.453502  
iteration 300 / 1500: loss 17.565047  
iteration 400 / 1500: loss 16.009198  
iteration 500 / 1500: loss 14.734748  
iteration 600 / 1500: loss 13.695182  
iteration 700 / 1500: loss 12.845738  
iteration 800 / 1500: loss 12.145728  
iteration 900 / 1500: loss 11.574274  
iteration 1000 / 1500: loss 11.106844  
iteration 1100 / 1500: loss 10.726786  
iteration 1200 / 1500: loss 10.413845  
iteration 1300 / 1500: loss 10.156120  
iteration 1400 / 1500: loss 9.942285  
iteration 0 / 1500: loss 41.215168  
iteration 100 / 1500: loss 30.586929  
iteration 200 / 1500: loss 23.459343  
iteration 300 / 1500: loss 18.687359  
iteration 400 / 1500: loss 15.492110  
iteration 500 / 1500: loss 13.353427  
iteration 600 / 1500: loss 11.912268  
iteration 700 / 1500: loss 10.955144  
iteration 800 / 1500: loss 10.307624

iteration 900 / 1500: loss 9.876250  
iteration 1000 / 1500: loss 9.587906  
iteration 1100 / 1500: loss 9.392915  
iteration 1200 / 1500: loss 9.264058  
iteration 1300 / 1500: loss 9.175502  
iteration 1400 / 1500: loss 9.116656  
iteration 0 / 1500: loss 56.992314  
iteration 100 / 1500: loss 35.318036  
iteration 200 / 1500: loss 23.428484  
iteration 300 / 1500: loss 16.912829  
iteration 400 / 1500: loss 13.338948  
iteration 500 / 1500: loss 11.377607  
iteration 600 / 1500: loss 10.304441  
iteration 700 / 1500: loss 9.713867  
iteration 800 / 1500: loss 9.392074  
iteration 900 / 1500: loss 9.214761  
iteration 1000 / 1500: loss 9.117002  
iteration 1100 / 1500: loss 9.064144  
iteration 1200 / 1500: loss 9.034676  
iteration 1300 / 1500: loss 9.018974  
iteration 1400 / 1500: loss 9.009795  
iteration 0 / 1500: loss 71.806144  
iteration 100 / 1500: loss 37.177618  
iteration 200 / 1500: loss 21.640609  
iteration 300 / 1500: loss 14.671909  
iteration 400 / 1500: loss 11.544970  
iteration 500 / 1500: loss 10.142297  
iteration 600 / 1500: loss 9.512133  
iteration 700 / 1500: loss 9.229340  
iteration 800 / 1500: loss 9.102448  
iteration 900 / 1500: loss 9.045769  
iteration 1000 / 1500: loss 9.020374  
iteration 1100 / 1500: loss 9.009014  
iteration 1200 / 1500: loss 9.003781  
iteration 1300 / 1500: loss 9.001471  
iteration 1400 / 1500: loss 9.000446  
iteration 0 / 1500: loss 88.089470  
iteration 100 / 1500: loss 38.027785  
iteration 200 / 1500: loss 19.648601  
iteration 300 / 1500: loss 12.908922  
iteration 400 / 1500: loss 10.437089  
iteration 500 / 1500: loss 9.525816  
iteration 600 / 1500: loss 9.192766  
iteration 700 / 1500: loss 9.070484  
iteration 800 / 1500: loss 9.025501  
iteration 900 / 1500: loss 9.009187  
iteration 1000 / 1500: loss 9.003053  
iteration 1100 / 1500: loss 9.000979  
iteration 1200 / 1500: loss 9.000162  
iteration 1300 / 1500: loss 8.999800  
iteration 1400 / 1500: loss 8.999701  
iteration 0 / 1500: loss 171.061714  
iteration 100 / 1500: loss 30.716381



iteration 200 / 1500: loss 11.909215  
iteration 300 / 1500: loss 9.389579  
iteration 400 / 1500: loss 9.052185  
iteration 500 / 1500: loss 9.006856  
iteration 600 / 1500: loss 9.000763  
iteration 700 / 1500: loss 8.999907  
iteration 800 / 1500: loss 8.999812  
iteration 900 / 1500: loss 8.999824  
iteration 1000 / 1500: loss 8.999816  
iteration 1100 / 1500: loss 8.999858  
iteration 1200 / 1500: loss 8.999816  
iteration 1300 / 1500: loss 8.999841  
iteration 1400 / 1500: loss 8.999859  
iteration 0 / 1500: loss 10.649769  
iteration 100 / 1500: loss 10.601363  
iteration 200 / 1500: loss 10.510027  
iteration 300 / 1500: loss 10.454438  
iteration 400 / 1500: loss 10.401835  
iteration 500 / 1500: loss 10.344445  
iteration 600 / 1500: loss 10.300307  
iteration 700 / 1500: loss 10.260208  
iteration 800 / 1500: loss 10.193068  
iteration 900 / 1500: loss 10.161515  
iteration 1000 / 1500: loss 10.106179  
iteration 1100 / 1500: loss 10.050118  
iteration 1200 / 1500: loss 10.007931  
iteration 1300 / 1500: loss 9.982270  
iteration 1400 / 1500: loss 9.938138  
iteration 0 / 1500: loss 17.387181  
iteration 100 / 1500: loss 15.866039  
iteration 200 / 1500: loss 14.617982  
iteration 300 / 1500: loss 13.600652  
iteration 400 / 1500: loss 12.771689  
iteration 500 / 1500: loss 12.079819  
iteration 600 / 1500: loss 11.524240  
iteration 700 / 1500: loss 11.064045  
iteration 800 / 1500: loss 10.686178  
iteration 900 / 1500: loss 10.379494  
iteration 1000 / 1500: loss 10.134963  
iteration 1100 / 1500: loss 9.930588  
iteration 1200 / 1500: loss 9.757500  
iteration 1300 / 1500: loss 9.618199  
iteration 1400 / 1500: loss 9.504882  
iteration 0 / 1500: loss 24.795060  
iteration 100 / 1500: loss 19.577447  
iteration 200 / 1500: loss 16.091228  
iteration 300 / 1500: loss 13.756382  
iteration 400 / 1500: loss 12.181977  
iteration 500 / 1500: loss 11.134599  
iteration 600 / 1500: loss 10.432101  
iteration 700 / 1500: loss 9.959688  
iteration 800 / 1500: loss 9.640130  
iteration 900 / 1500: loss 9.429855

iteration 1000 / 1500: loss 9.287308  
iteration 1100 / 1500: loss 9.190578  
iteration 1200 / 1500: loss 9.125790  
iteration 1300 / 1500: loss 9.084592  
iteration 1400 / 1500: loss 9.056440  
iteration 0 / 1500: loss 41.984678  
iteration 100 / 1500: loss 23.792966  
iteration 200 / 1500: loss 15.632379  
iteration 300 / 1500: loss 11.978102  
iteration 400 / 1500: loss 10.336052  
iteration 500 / 1500: loss 9.596862  
iteration 600 / 1500: loss 9.268282  
iteration 700 / 1500: loss 9.120058  
iteration 800 / 1500: loss 9.053599  
iteration 900 / 1500: loss 9.022866  
iteration 1000 / 1500: loss 9.009820  
iteration 1100 / 1500: loss 9.003661  
iteration 1200 / 1500: loss 9.001228  
iteration 1300 / 1500: loss 9.000169  
iteration 1400 / 1500: loss 8.999427  
iteration 0 / 1500: loss 54.622317  
iteration 100 / 1500: loss 22.694929  
iteration 200 / 1500: loss 13.106683  
iteration 300 / 1500: loss 10.235222  
iteration 400 / 1500: loss 9.369835  
iteration 500 / 1500: loss 9.110756  
iteration 600 / 1500: loss 9.032846  
iteration 700 / 1500: loss 9.009534  
iteration 800 / 1500: loss 9.002205  
iteration 900 / 1500: loss 9.000426  
iteration 1000 / 1500: loss 8.999723  
iteration 1100 / 1500: loss 8.999557  
iteration 1200 / 1500: loss 8.999541  
iteration 1300 / 1500: loss 8.999424  
iteration 1400 / 1500: loss 8.999423  
iteration 0 / 1500: loss 73.772330  
iteration 100 / 1500: loss 21.988922  
iteration 200 / 1500: loss 11.608119  
iteration 300 / 1500: loss 9.521271  
iteration 400 / 1500: loss 9.104405  
iteration 500 / 1500: loss 9.020349  
iteration 600 / 1500: loss 9.003797  
iteration 700 / 1500: loss 9.000310  
iteration 800 / 1500: loss 8.999705  
iteration 900 / 1500: loss 8.999640  
iteration 1000 / 1500: loss 8.999554  
iteration 1100 / 1500: loss 8.999570  
iteration 1200 / 1500: loss 8.999551  
iteration 1300 / 1500: loss 8.999696  
iteration 1400 / 1500: loss 8.999587  
iteration 0 / 1500: loss 87.630900  
iteration 100 / 1500: loss 19.528651  
iteration 200 / 1500: loss 10.410475

iteration 300 / 1500: loss 9.188930  
iteration 400 / 1500: loss 9.024899  
iteration 500 / 1500: loss 9.003082  
iteration 600 / 1500: loss 9.000108  
iteration 700 / 1500: loss 8.999695  
iteration 800 / 1500: loss 8.999671  
iteration 900 / 1500: loss 8.999654  
iteration 1000 / 1500: loss 8.999692  
iteration 1100 / 1500: loss 8.999623  
iteration 1200 / 1500: loss 8.999664  
iteration 1300 / 1500: loss 8.999582  
iteration 1400 / 1500: loss 8.999663  
iteration 0 / 1500: loss 169.096718  
iteration 100 / 1500: loss 11.814801  
iteration 200 / 1500: loss 9.049346  
iteration 300 / 1500: loss 9.000670  
iteration 400 / 1500: loss 8.999870  
iteration 500 / 1500: loss 8.999839  
iteration 600 / 1500: loss 8.999842  
iteration 700 / 1500: loss 8.999862  
iteration 800 / 1500: loss 8.999791  
iteration 900 / 1500: loss 8.999797  
iteration 1000 / 1500: loss 8.999793  
iteration 1100 / 1500: loss 8.999786  
iteration 1200 / 1500: loss 8.999813  
iteration 1300 / 1500: loss 8.999809  
iteration 1400 / 1500: loss 8.999826  
iteration 0 / 1500: loss 10.604011  
iteration 100 / 1500: loss 10.472676  
iteration 200 / 1500: loss 10.364896  
iteration 300 / 1500: loss 10.261710  
iteration 400 / 1500: loss 10.167433  
iteration 500 / 1500: loss 10.074861  
iteration 600 / 1500: loss 9.979947  
iteration 700 / 1500: loss 9.914838  
iteration 800 / 1500: loss 9.828529  
iteration 900 / 1500: loss 9.781575  
iteration 1000 / 1500: loss 9.710182  
iteration 1100 / 1500: loss 9.670637  
iteration 1200 / 1500: loss 9.611616  
iteration 1300 / 1500: loss 9.563630  
iteration 1400 / 1500: loss 9.509502  
iteration 0 / 1500: loss 16.966789  
iteration 100 / 1500: loss 14.350894  
iteration 200 / 1500: loss 12.579941  
iteration 300 / 1500: loss 11.394810  
iteration 400 / 1500: loss 10.613195  
iteration 500 / 1500: loss 10.076459  
iteration 600 / 1500: loss 9.720416  
iteration 700 / 1500: loss 9.480665  
iteration 800 / 1500: loss 9.319702  
iteration 900 / 1500: loss 9.213474  
iteration 1000 / 1500: loss 9.142741

iteration 1100 / 1500: loss 9.094373  
iteration 1200 / 1500: loss 9.061986  
iteration 1300 / 1500: loss 9.040417  
iteration 1400 / 1500: loss 9.025286  
iteration 0 / 1500: loss 25.081834  
iteration 100 / 1500: loss 16.212045  
iteration 200 / 1500: loss 12.230897  
iteration 300 / 1500: loss 10.446150  
iteration 400 / 1500: loss 9.647403  
iteration 500 / 1500: loss 9.290825  
iteration 600 / 1500: loss 9.128026  
iteration 700 / 1500: loss 9.057532  
iteration 800 / 1500: loss 9.024519  
iteration 900 / 1500: loss 9.010335  
iteration 1000 / 1500: loss 9.003811  
iteration 1100 / 1500: loss 9.000578  
iteration 1200 / 1500: loss 8.999315  
iteration 1300 / 1500: loss 8.998705  
iteration 1400 / 1500: loss 8.998336  
iteration 0 / 1500: loss 41.144295  
iteration 100 / 1500: loss 15.447825  
iteration 200 / 1500: loss 10.292492  
iteration 300 / 1500: loss 9.259019  
iteration 400 / 1500: loss 9.051028  
iteration 500 / 1500: loss 9.009489  
iteration 600 / 1500: loss 9.001051  
iteration 700 / 1500: loss 8.999504  
iteration 800 / 1500: loss 8.999237  
iteration 900 / 1500: loss 8.999150  
iteration 1000 / 1500: loss 8.999052  
iteration 1100 / 1500: loss 8.999022  
iteration 1200 / 1500: loss 8.999202  
iteration 1300 / 1500: loss 8.999341  
iteration 1400 / 1500: loss 8.999178  
iteration 0 / 1500: loss 58.600333  
iteration 100 / 1500: loss 13.433920  
iteration 200 / 1500: loss 9.395946  
iteration 300 / 1500: loss 9.034582  
iteration 400 / 1500: loss 9.002482  
iteration 500 / 1500: loss 8.999674  
iteration 600 / 1500: loss 8.999379  
iteration 700 / 1500: loss 8.999499  
iteration 800 / 1500: loss 8.999400  
iteration 900 / 1500: loss 8.999474  
iteration 1000 / 1500: loss 8.999220  
iteration 1100 / 1500: loss 8.999395  
iteration 1200 / 1500: loss 8.999348  
iteration 1300 / 1500: loss 8.999499  
iteration 1400 / 1500: loss 8.999431  
iteration 0 / 1500: loss 75.681033  
iteration 100 / 1500: loss 11.651563  
iteration 200 / 1500: loss 9.104614  
iteration 300 / 1500: loss 9.003786

iteration 400 / 1500: loss 8.999635  
iteration 500 / 1500: loss 8.999590  
iteration 600 / 1500: loss 8.999578  
iteration 700 / 1500: loss 8.999654  
iteration 800 / 1500: loss 8.999546  
iteration 900 / 1500: loss 8.999539  
iteration 1000 / 1500: loss 8.999626  
iteration 1100 / 1500: loss 8.999588  
iteration 1200 / 1500: loss 8.999535  
iteration 1300 / 1500: loss 8.999426  
iteration 1400 / 1500: loss 8.999521  
iteration 0 / 1500: loss 90.584749  
iteration 100 / 1500: loss 10.432575  
iteration 200 / 1500: loss 9.025119  
iteration 300 / 1500: loss 9.000137  
iteration 400 / 1500: loss 8.999540  
iteration 500 / 1500: loss 8.999557  
iteration 600 / 1500: loss 8.999643  
iteration 700 / 1500: loss 8.999586  
iteration 800 / 1500: loss 8.999671  
iteration 900 / 1500: loss 8.999660  
iteration 1000 / 1500: loss 8.999664  
iteration 1100 / 1500: loss 8.999622  
iteration 1200 / 1500: loss 8.999616  
iteration 1300 / 1500: loss 8.999588  
iteration 1400 / 1500: loss 8.999640  
iteration 0 / 1500: loss 173.694308  
iteration 100 / 1500: loss 9.046457  
iteration 200 / 1500: loss 8.999853  
iteration 300 / 1500: loss 8.999820  
iteration 400 / 1500: loss 8.999819  
iteration 500 / 1500: loss 8.999787  
iteration 600 / 1500: loss 8.999841  
iteration 700 / 1500: loss 8.999827  
iteration 800 / 1500: loss 8.999826  
iteration 900 / 1500: loss 8.999828  
iteration 1000 / 1500: loss 8.999815  
iteration 1100 / 1500: loss 8.999838  
iteration 1200 / 1500: loss 8.999837  
iteration 1300 / 1500: loss 8.999824  
iteration 1400 / 1500: loss 8.999834  
iteration 0 / 1500: loss 10.597976  
iteration 100 / 1500: loss 10.422787  
iteration 200 / 1500: loss 10.261017  
iteration 300 / 1500: loss 10.116479  
iteration 400 / 1500: loss 9.975429  
iteration 500 / 1500: loss 9.873461  
iteration 600 / 1500: loss 9.773045  
iteration 700 / 1500: loss 9.676795  
iteration 800 / 1500: loss 9.597652  
iteration 900 / 1500: loss 9.521525  
iteration 1000 / 1500: loss 9.463032  
iteration 1100 / 1500: loss 9.420537

iteration 1200 / 1500: loss 9.366744  
iteration 1300 / 1500: loss 9.323150  
iteration 1400 / 1500: loss 9.282847  
iteration 0 / 1500: loss 17.233463  
iteration 100 / 1500: loss 13.505898  
iteration 200 / 1500: loss 11.479238  
iteration 300 / 1500: loss 10.357619  
iteration 400 / 1500: loss 9.741784  
iteration 500 / 1500: loss 9.404683  
iteration 600 / 1500: loss 9.222597  
iteration 700 / 1500: loss 9.120459  
iteration 800 / 1500: loss 9.062282  
iteration 900 / 1500: loss 9.032998  
iteration 1000 / 1500: loss 9.016069  
iteration 1100 / 1500: loss 9.007595  
iteration 1200 / 1500: loss 9.002323  
iteration 1300 / 1500: loss 8.999683  
iteration 1400 / 1500: loss 8.998030  
iteration 0 / 1500: loss 25.399315  
iteration 100 / 1500: loss 13.916744  
iteration 200 / 1500: loss 10.472596  
iteration 300 / 1500: loss 9.441882  
iteration 400 / 1500: loss 9.129961  
iteration 500 / 1500: loss 9.038050  
iteration 600 / 1500: loss 9.010227  
iteration 700 / 1500: loss 9.001496  
iteration 800 / 1500: loss 8.999173  
iteration 900 / 1500: loss 8.998699  
iteration 1000 / 1500: loss 8.997935  
iteration 1100 / 1500: loss 8.998308  
iteration 1200 / 1500: loss 8.998244  
iteration 1300 / 1500: loss 8.998004  
iteration 1400 / 1500: loss 8.998457  
iteration 0 / 1500: loss 41.763473  
iteration 100 / 1500: loss 11.928966  
iteration 200 / 1500: loss 9.260535  
iteration 300 / 1500: loss 9.022207  
iteration 400 / 1500: loss 9.001227  
iteration 500 / 1500: loss 8.999098  
iteration 600 / 1500: loss 8.999064  
iteration 700 / 1500: loss 8.999171  
iteration 800 / 1500: loss 8.999094  
iteration 900 / 1500: loss 8.999110  
iteration 1000 / 1500: loss 8.999224  
iteration 1100 / 1500: loss 8.999188  
iteration 1200 / 1500: loss 8.999206  
iteration 1300 / 1500: loss 8.999165  
iteration 1400 / 1500: loss 8.998876  
iteration 0 / 1500: loss 57.657941  
iteration 100 / 1500: loss 10.287605  
iteration 200 / 1500: loss 9.033523  
iteration 300 / 1500: loss 9.000453  
iteration 400 / 1500: loss 8.999401

iteration 500 / 1500: loss 8.999362  
iteration 600 / 1500: loss 8.999563  
iteration 700 / 1500: loss 8.999465  
iteration 800 / 1500: loss 8.999545  
iteration 900 / 1500: loss 8.999398  
iteration 1000 / 1500: loss 8.999354  
iteration 1100 / 1500: loss 8.999484  
iteration 1200 / 1500: loss 8.999383  
iteration 1300 / 1500: loss 8.999403  
iteration 1400 / 1500: loss 8.999410  
iteration 0 / 1500: loss 75.892270  
iteration 100 / 1500: loss 9.520302  
iteration 200 / 1500: loss 9.003729  
iteration 300 / 1500: loss 8.999623  
iteration 400 / 1500: loss 8.999502  
iteration 500 / 1500: loss 8.999656  
iteration 600 / 1500: loss 8.999631  
iteration 700 / 1500: loss 8.999546  
iteration 800 / 1500: loss 8.999476  
iteration 900 / 1500: loss 8.999507  
iteration 1000 / 1500: loss 8.999578  
iteration 1100 / 1500: loss 8.999559  
iteration 1200 / 1500: loss 8.999521  
iteration 1300 / 1500: loss 8.999553  
iteration 1400 / 1500: loss 8.999527  
iteration 0 / 1500: loss 89.518852  
iteration 100 / 1500: loss 9.181799  
iteration 200 / 1500: loss 9.000124  
iteration 300 / 1500: loss 8.999649  
iteration 400 / 1500: loss 8.999701  
iteration 500 / 1500: loss 8.999619  
iteration 600 / 1500: loss 8.999628  
iteration 700 / 1500: loss 8.999640  
iteration 800 / 1500: loss 8.999594  
iteration 900 / 1500: loss 8.999611  
iteration 1000 / 1500: loss 8.999734  
iteration 1100 / 1500: loss 8.999569  
iteration 1200 / 1500: loss 8.999629  
iteration 1300 / 1500: loss 8.999596  
iteration 1400 / 1500: loss 8.999583  
iteration 0 / 1500: loss 161.811135  
iteration 100 / 1500: loss 9.000467  
iteration 200 / 1500: loss 8.999852  
iteration 300 / 1500: loss 8.999852  
iteration 400 / 1500: loss 8.999797  
iteration 500 / 1500: loss 8.999783  
iteration 600 / 1500: loss 8.999849  
iteration 700 / 1500: loss 8.999827  
iteration 800 / 1500: loss 8.999826  
iteration 900 / 1500: loss 8.999832  
iteration 1000 / 1500: loss 8.999842  
iteration 1100 / 1500: loss 8.999863  
iteration 1200 / 1500: loss 8.999851

iteration 1300 / 1500: loss 8.999799  
iteration 1400 / 1500: loss 8.999831  
iteration 0 / 1500: loss 10.729373  
iteration 100 / 1500: loss 10.465829  
iteration 200 / 1500: loss 10.264722  
iteration 300 / 1500: loss 10.064632  
iteration 400 / 1500: loss 9.909230  
iteration 500 / 1500: loss 9.758948  
iteration 600 / 1500: loss 9.651990  
iteration 700 / 1500: loss 9.559395  
iteration 800 / 1500: loss 9.472006  
iteration 900 / 1500: loss 9.403993  
iteration 1000 / 1500: loss 9.333529  
iteration 1100 / 1500: loss 9.281952  
iteration 1200 / 1500: loss 9.235019  
iteration 1300 / 1500: loss 9.203810  
iteration 1400 / 1500: loss 9.165482  
iteration 0 / 1500: loss 17.327822  
iteration 100 / 1500: loss 12.731024  
iteration 200 / 1500: loss 10.664237  
iteration 300 / 1500: loss 9.747534  
iteration 400 / 1500: loss 9.337824  
iteration 500 / 1500: loss 9.147404  
iteration 600 / 1500: loss 9.063699  
iteration 700 / 1500: loss 9.027401  
iteration 800 / 1500: loss 9.010290  
iteration 900 / 1500: loss 9.002423  
iteration 1000 / 1500: loss 8.998465  
iteration 1100 / 1500: loss 8.996756  
iteration 1200 / 1500: loss 8.996777  
iteration 1300 / 1500: loss 8.996360  
iteration 1400 / 1500: loss 8.996326  
iteration 0 / 1500: loss 24.396710  
iteration 100 / 1500: loss 12.090010  
iteration 200 / 1500: loss 9.618799  
iteration 300 / 1500: loss 9.122673  
iteration 400 / 1500: loss 9.023045  
iteration 500 / 1500: loss 9.003485  
iteration 600 / 1500: loss 8.999619  
iteration 700 / 1500: loss 8.998247  
iteration 800 / 1500: loss 8.998359  
iteration 900 / 1500: loss 8.997749  
iteration 1000 / 1500: loss 8.998128  
iteration 1100 / 1500: loss 8.998368  
iteration 1200 / 1500: loss 8.998521  
iteration 1300 / 1500: loss 8.998490  
iteration 1400 / 1500: loss 8.998483  
iteration 0 / 1500: loss 42.884748  
iteration 100 / 1500: loss 10.346457  
iteration 200 / 1500: loss 9.052748  
iteration 300 / 1500: loss 9.001272  
iteration 400 / 1500: loss 8.999069  
iteration 500 / 1500: loss 8.998997



iteration 600 / 1500: loss 8.999062  
iteration 700 / 1500: loss 8.999167  
iteration 800 / 1500: loss 8.999007  
iteration 900 / 1500: loss 8.999224  
iteration 1000 / 1500: loss 8.999115  
iteration 1100 / 1500: loss 8.999282  
iteration 1200 / 1500: loss 8.999243  
iteration 1300 / 1500: loss 8.999141  
iteration 1400 / 1500: loss 8.999122  
iteration 0 / 1500: loss 57.042659  
iteration 100 / 1500: loss 9.372103  
iteration 200 / 1500: loss 9.002303  
iteration 300 / 1500: loss 8.999517  
iteration 400 / 1500: loss 8.999503  
iteration 500 / 1500: loss 8.999315  
iteration 600 / 1500: loss 8.999387  
iteration 700 / 1500: loss 8.999357  
iteration 800 / 1500: loss 8.999399  
iteration 900 / 1500: loss 8.999391  
iteration 1000 / 1500: loss 8.999494  
iteration 1100 / 1500: loss 8.999353  
iteration 1200 / 1500: loss 8.999349  
iteration 1300 / 1500: loss 8.999480  
iteration 1400 / 1500: loss 8.999298  
iteration 0 / 1500: loss 72.398296  
iteration 100 / 1500: loss 9.094401  
iteration 200 / 1500: loss 8.999673  
iteration 300 / 1500: loss 8.999496  
iteration 400 / 1500: loss 8.999612  
iteration 500 / 1500: loss 8.999604  
iteration 600 / 1500: loss 8.999550  
iteration 700 / 1500: loss 8.999600  
iteration 800 / 1500: loss 8.999535  
iteration 900 / 1500: loss 8.999512  
iteration 1000 / 1500: loss 8.999676  
iteration 1100 / 1500: loss 8.999496  
iteration 1200 / 1500: loss 8.999633  
iteration 1300 / 1500: loss 8.999523  
iteration 1400 / 1500: loss 8.999590  
iteration 0 / 1500: loss 90.342942  
iteration 100 / 1500: loss 9.022753  
iteration 200 / 1500: loss 8.999724  
iteration 300 / 1500: loss 8.999726  
iteration 400 / 1500: loss 8.999644  
iteration 500 / 1500: loss 8.999647  
iteration 600 / 1500: loss 8.999721  
iteration 700 / 1500: loss 8.999716  
iteration 800 / 1500: loss 8.999638  
iteration 900 / 1500: loss 8.999684  
iteration 1000 / 1500: loss 8.999610  
iteration 1100 / 1500: loss 8.999708  
iteration 1200 / 1500: loss 8.999627  
iteration 1300 / 1500: loss 8.999707

iteration 1400 / 1500: loss 8.999655  
iteration 0 / 1500: loss 173.556081  
iteration 100 / 1500: loss 8.999862  
iteration 200 / 1500: loss 8.999848  
iteration 300 / 1500: loss 8.999832  
iteration 400 / 1500: loss 8.999807  
iteration 500 / 1500: loss 8.999828  
iteration 600 / 1500: loss 8.999817  
iteration 700 / 1500: loss 8.999855  
iteration 800 / 1500: loss 8.999844  
iteration 900 / 1500: loss 8.999846  
iteration 1000 / 1500: loss 8.999830  
iteration 1100 / 1500: loss 8.999879  
iteration 1200 / 1500: loss 8.999808  
iteration 1300 / 1500: loss 8.999812  
iteration 1400 / 1500: loss 8.999814  
iteration 0 / 1500: loss 10.570393  
iteration 100 / 1500: loss 10.264311  
iteration 200 / 1500: loss 10.047732  
iteration 300 / 1500: loss 9.855832  
iteration 400 / 1500: loss 9.684709  
iteration 500 / 1500: loss 9.565026  
iteration 600 / 1500: loss 9.451793  
iteration 700 / 1500: loss 9.367290  
iteration 800 / 1500: loss 9.297359  
iteration 900 / 1500: loss 9.247584  
iteration 1000 / 1500: loss 9.196811  
iteration 1100 / 1500: loss 9.159311  
iteration 1200 / 1500: loss 9.122963  
iteration 1300 / 1500: loss 9.105164  
iteration 1400 / 1500: loss 9.079747  
iteration 0 / 1500: loss 16.981030  
iteration 100 / 1500: loss 11.920799  
iteration 200 / 1500: loss 10.069927  
iteration 300 / 1500: loss 9.392087  
iteration 400 / 1500: loss 9.143255  
iteration 500 / 1500: loss 9.048390  
iteration 600 / 1500: loss 9.015748  
iteration 700 / 1500: loss 9.004352  
iteration 800 / 1500: loss 8.999161  
iteration 900 / 1500: loss 8.997283  
iteration 1000 / 1500: loss 8.997040  
iteration 1100 / 1500: loss 8.996720  
iteration 1200 / 1500: loss 8.996523  
iteration 1300 / 1500: loss 8.997189  
iteration 1400 / 1500: loss 8.996433  
iteration 0 / 1500: loss 25.307496  
iteration 100 / 1500: loss 11.188444  
iteration 200 / 1500: loss 9.293045  
iteration 300 / 1500: loss 9.037573  
iteration 400 / 1500: loss 9.003320  
iteration 500 / 1500: loss 8.999025  
iteration 600 / 1500: loss 8.998391

iteration 700 / 1500: loss 8.998195  
iteration 800 / 1500: loss 8.998018  
iteration 900 / 1500: loss 8.998440  
iteration 1000 / 1500: loss 8.998252  
iteration 1100 / 1500: loss 8.998526  
iteration 1200 / 1500: loss 8.998471  
iteration 1300 / 1500: loss 8.997971  
iteration 1400 / 1500: loss 8.998179  
iteration 0 / 1500: loss 41.368160  
iteration 100 / 1500: loss 9.569448  
iteration 200 / 1500: loss 9.008812  
iteration 300 / 1500: loss 8.999039  
iteration 400 / 1500: loss 8.999124  
iteration 500 / 1500: loss 8.999072  
iteration 600 / 1500: loss 8.999272  
iteration 700 / 1500: loss 8.998984  
iteration 800 / 1500: loss 8.999159  
iteration 900 / 1500: loss 8.998962  
iteration 1000 / 1500: loss 8.999010  
iteration 1100 / 1500: loss 8.999078  
iteration 1200 / 1500: loss 8.999084  
iteration 1300 / 1500: loss 8.999181  
iteration 1400 / 1500: loss 8.999083  
iteration 0 / 1500: loss 54.815418  
iteration 100 / 1500: loss 9.102831  
iteration 200 / 1500: loss 8.999669  
iteration 300 / 1500: loss 8.999447  
iteration 400 / 1500: loss 8.999405  
iteration 500 / 1500: loss 8.999452  
iteration 600 / 1500: loss 8.999530  
iteration 700 / 1500: loss 8.999408  
iteration 800 / 1500: loss 8.999338  
iteration 900 / 1500: loss 8.999496  
iteration 1000 / 1500: loss 8.999509  
iteration 1100 / 1500: loss 8.999203  
iteration 1200 / 1500: loss 8.999437  
iteration 1300 / 1500: loss 8.999345  
iteration 1400 / 1500: loss 8.999330  
iteration 0 / 1500: loss 75.362154  
iteration 100 / 1500: loss 9.018300  
iteration 200 / 1500: loss 8.999580  
iteration 300 / 1500: loss 8.999650  
iteration 400 / 1500: loss 8.999458  
iteration 500 / 1500: loss 8.999642  
iteration 600 / 1500: loss 8.999487  
iteration 700 / 1500: loss 8.999515  
iteration 800 / 1500: loss 8.999649  
iteration 900 / 1500: loss 8.999542  
iteration 1000 / 1500: loss 8.999501  
iteration 1100 / 1500: loss 8.999602  
iteration 1200 / 1500: loss 8.999582  
iteration 1300 / 1500: loss 8.999585  
iteration 1400 / 1500: loss 8.999560

iteration 0 / 1500: loss 86.802750  
iteration 100 / 1500: loss 9.002371  
iteration 200 / 1500: loss 8.999649  
iteration 300 / 1500: loss 8.999634  
iteration 400 / 1500: loss 8.999653  
iteration 500 / 1500: loss 8.999657  
iteration 600 / 1500: loss 8.999642  
iteration 700 / 1500: loss 8.999632  
iteration 800 / 1500: loss 8.999622  
iteration 900 / 1500: loss 8.999620  
iteration 1000 / 1500: loss 8.999635  
iteration 1100 / 1500: loss 8.999676  
iteration 1200 / 1500: loss 8.999669  
iteration 1300 / 1500: loss 8.999626  
iteration 1400 / 1500: loss 8.999655  
iteration 0 / 1500: loss 161.531972  
iteration 100 / 1500: loss 8.999815  
iteration 200 / 1500: loss 8.999822  
iteration 300 / 1500: loss 8.999835  
iteration 400 / 1500: loss 8.999802  
iteration 500 / 1500: loss 8.999814  
iteration 600 / 1500: loss 8.999784  
iteration 700 / 1500: loss 8.999858  
iteration 800 / 1500: loss 8.999833  
iteration 900 / 1500: loss 8.999802  
iteration 1000 / 1500: loss 8.999793  
iteration 1100 / 1500: loss 8.999823  
iteration 1200 / 1500: loss 8.999889  
iteration 1300 / 1500: loss 8.999873  
iteration 1400 / 1500: loss 8.999832  
iteration 0 / 1500: loss 10.587828  
iteration 100 / 1500: loss 10.052891  
iteration 200 / 1500: loss 9.702093  
iteration 300 / 1500: loss 9.458461  
iteration 400 / 1500: loss 9.306641  
iteration 500 / 1500: loss 9.196640  
iteration 600 / 1500: loss 9.128164  
iteration 700 / 1500: loss 9.076964  
iteration 800 / 1500: loss 9.044735  
iteration 900 / 1500: loss 9.026516  
iteration 1000 / 1500: loss 9.010755  
iteration 1100 / 1500: loss 9.001037  
iteration 1200 / 1500: loss 8.994874  
iteration 1300 / 1500: loss 8.989765  
iteration 1400 / 1500: loss 8.989583  
iteration 0 / 1500: loss 16.568270  
iteration 100 / 1500: loss 10.011221  
iteration 200 / 1500: loss 9.133680  
iteration 300 / 1500: loss 9.013916  
iteration 400 / 1500: loss 8.998662  
iteration 500 / 1500: loss 8.996388  
iteration 600 / 1500: loss 8.996028  
iteration 700 / 1500: loss 8.996395

iteration 800 / 1500: loss 8.996658  
iteration 900 / 1500: loss 8.995615  
iteration 1000 / 1500: loss 8.996330  
iteration 1100 / 1500: loss 8.996738  
iteration 1200 / 1500: loss 8.996399  
iteration 1300 / 1500: loss 8.996673  
iteration 1400 / 1500: loss 8.996325  
iteration 0 / 1500: loss 24.854267  
iteration 100 / 1500: loss 9.278244  
iteration 200 / 1500: loss 9.002978  
iteration 300 / 1500: loss 8.998335  
iteration 400 / 1500: loss 8.998720  
iteration 500 / 1500: loss 8.998517  
iteration 600 / 1500: loss 8.997834  
iteration 700 / 1500: loss 8.998193  
iteration 800 / 1500: loss 8.997918  
iteration 900 / 1500: loss 8.998207  
iteration 1000 / 1500: loss 8.998369  
iteration 1100 / 1500: loss 8.997978  
iteration 1200 / 1500: loss 8.998069  
iteration 1300 / 1500: loss 8.998181  
iteration 1400 / 1500: loss 8.998478  
iteration 0 / 1500: loss 41.654656  
iteration 100 / 1500: loss 9.008239  
iteration 200 / 1500: loss 8.998992  
iteration 300 / 1500: loss 8.998919  
iteration 400 / 1500: loss 8.999109  
iteration 500 / 1500: loss 8.999208  
iteration 600 / 1500: loss 8.999118  
iteration 700 / 1500: loss 8.998965  
iteration 800 / 1500: loss 8.999188  
iteration 900 / 1500: loss 8.998938  
iteration 1000 / 1500: loss 8.999082  
iteration 1100 / 1500: loss 8.999238  
iteration 1200 / 1500: loss 8.999082  
iteration 1300 / 1500: loss 8.999304  
iteration 1400 / 1500: loss 8.999151  
iteration 0 / 1500: loss 58.108084  
iteration 100 / 1500: loss 8.999538  
iteration 200 / 1500: loss 8.999393  
iteration 300 / 1500: loss 8.999497  
iteration 400 / 1500: loss 8.999308  
iteration 500 / 1500: loss 8.999375  
iteration 600 / 1500: loss 8.999478  
iteration 700 / 1500: loss 8.999545  
iteration 800 / 1500: loss 8.999515  
iteration 900 / 1500: loss 8.999319  
iteration 1000 / 1500: loss 8.999495  
iteration 1100 / 1500: loss 8.999382  
iteration 1200 / 1500: loss 8.999423  
iteration 1300 / 1500: loss 8.999396  
iteration 1400 / 1500: loss 8.999458  
iteration 0 / 1500: loss 75.523520

iteration 100 / 1500: loss 8.999494  
iteration 200 / 1500: loss 8.999553  
iteration 300 / 1500: loss 8.999547  
iteration 400 / 1500: loss 8.999507  
iteration 500 / 1500: loss 8.999614  
iteration 600 / 1500: loss 8.999510  
iteration 700 / 1500: loss 8.999666  
iteration 800 / 1500: loss 8.999523  
iteration 900 / 1500: loss 8.999602  
iteration 1000 / 1500: loss 8.999673  
iteration 1100 / 1500: loss 8.999549  
iteration 1200 / 1500: loss 8.999615  
iteration 1300 / 1500: loss 8.999557  
iteration 1400 / 1500: loss 8.999568  
iteration 0 / 1500: loss 85.872054  
iteration 100 / 1500: loss 8.999647  
iteration 200 / 1500: loss 8.999685  
iteration 300 / 1500: loss 8.999686  
iteration 400 / 1500: loss 8.999697  
iteration 500 / 1500: loss 8.999716  
iteration 600 / 1500: loss 8.999745  
iteration 700 / 1500: loss 8.999635  
iteration 800 / 1500: loss 8.999690  
iteration 900 / 1500: loss 8.999684  
iteration 1000 / 1500: loss 8.999669  
iteration 1100 / 1500: loss 8.999610  
iteration 1200 / 1500: loss 8.999605  
iteration 1300 / 1500: loss 8.999643  
iteration 1400 / 1500: loss 8.999704  
iteration 0 / 1500: loss 173.709683  
iteration 100 / 1500: loss 8.999843  
iteration 200 / 1500: loss 8.999824  
iteration 300 / 1500: loss 8.999836  
iteration 400 / 1500: loss 8.999850  
iteration 500 / 1500: loss 8.999813  
iteration 600 / 1500: loss 8.999864  
iteration 700 / 1500: loss 8.999886  
iteration 800 / 1500: loss 8.999812  
iteration 900 / 1500: loss 8.999853  
iteration 1000 / 1500: loss 8.999815  
iteration 1100 / 1500: loss 8.999854  
iteration 1200 / 1500: loss 8.999830  
iteration 1300 / 1500: loss 8.999827  
iteration 1400 / 1500: loss 8.999812  
iteration 0 / 1500: loss 10.644058  
iteration 100 / 1500: loss 9.207834  
iteration 200 / 1500: loss 9.011109  
iteration 300 / 1500: loss 8.988933  
iteration 400 / 1500: loss 8.984772  
iteration 500 / 1500: loss 8.983945  
iteration 600 / 1500: loss 8.983209  
iteration 700 / 1500: loss 8.980983  
iteration 800 / 1500: loss 8.981477

iteration 900 / 1500: loss 8.981896  
iteration 1000 / 1500: loss 8.979910  
iteration 1100 / 1500: loss 8.983869  
iteration 1200 / 1500: loss 8.979598  
iteration 1300 / 1500: loss 8.980929  
iteration 1400 / 1500: loss 8.983668  
iteration 0 / 1500: loss 17.379961  
iteration 100 / 1500: loss 8.996374  
iteration 200 / 1500: loss 8.995983  
iteration 300 / 1500: loss 8.995851  
iteration 400 / 1500: loss 8.997152  
iteration 500 / 1500: loss 8.996199  
iteration 600 / 1500: loss 8.996209  
iteration 700 / 1500: loss 8.995864  
iteration 800 / 1500: loss 8.995865  
iteration 900 / 1500: loss 8.996624  
iteration 1000 / 1500: loss 8.996295  
iteration 1100 / 1500: loss 8.995846  
iteration 1200 / 1500: loss 8.996987  
iteration 1300 / 1500: loss 8.996483  
iteration 1400 / 1500: loss 8.996271  
iteration 0 / 1500: loss 25.867173  
iteration 100 / 1500: loss 8.998509  
iteration 200 / 1500: loss 8.998708  
iteration 300 / 1500: loss 8.998210  
iteration 400 / 1500: loss 8.998283  
iteration 500 / 1500: loss 8.998032  
iteration 600 / 1500: loss 8.998144  
iteration 700 / 1500: loss 8.998215  
iteration 800 / 1500: loss 8.998542  
iteration 900 / 1500: loss 8.998573  
iteration 1000 / 1500: loss 8.997999  
iteration 1100 / 1500: loss 8.998693  
iteration 1200 / 1500: loss 8.998112  
iteration 1300 / 1500: loss 8.998575  
iteration 1400 / 1500: loss 8.997898  
iteration 0 / 1500: loss 42.014623  
iteration 100 / 1500: loss 8.999297  
iteration 200 / 1500: loss 8.998963  
iteration 300 / 1500: loss 8.999306  
iteration 400 / 1500: loss 8.999105  
iteration 500 / 1500: loss 8.999178  
iteration 600 / 1500: loss 8.999129  
iteration 700 / 1500: loss 8.999370  
iteration 800 / 1500: loss 8.999198  
iteration 900 / 1500: loss 8.999138  
iteration 1000 / 1500: loss 8.999369  
iteration 1100 / 1500: loss 8.999068  
iteration 1200 / 1500: loss 8.999063  
iteration 1300 / 1500: loss 8.999171  
iteration 1400 / 1500: loss 8.999220  
iteration 0 / 1500: loss 57.341200  
iteration 100 / 1500: loss 8.999549

iteration 200 / 1500: loss 8.999563  
iteration 300 / 1500: loss 8.999709  
iteration 400 / 1500: loss 8.999482  
iteration 500 / 1500: loss 8.999421  
iteration 600 / 1500: loss 8.999578  
iteration 700 / 1500: loss 8.999588  
iteration 800 / 1500: loss 8.999414  
iteration 900 / 1500: loss 8.999363  
iteration 1000 / 1500: loss 8.999493  
iteration 1100 / 1500: loss 8.999473  
iteration 1200 / 1500: loss 8.999486  
iteration 1300 / 1500: loss 8.999490  
iteration 1400 / 1500: loss 8.999658  
iteration 0 / 1500: loss 74.177490  
iteration 100 / 1500: loss 8.999791  
iteration 200 / 1500: loss 8.999668  
iteration 300 / 1500: loss 8.999729  
iteration 400 / 1500: loss 8.999738  
iteration 500 / 1500: loss 8.999627  
iteration 600 / 1500: loss 8.999656  
iteration 700 / 1500: loss 8.999694  
iteration 800 / 1500: loss 8.999696  
iteration 900 / 1500: loss 8.999695  
iteration 1000 / 1500: loss 8.999639  
iteration 1100 / 1500: loss 8.999769  
iteration 1200 / 1500: loss 8.999578  
iteration 1300 / 1500: loss 8.999685  
iteration 1400 / 1500: loss 8.999714  
iteration 0 / 1500: loss 93.935979  
iteration 100 / 1500: loss 8.999766  
iteration 200 / 1500: loss 8.999712  
iteration 300 / 1500: loss 8.999835  
iteration 400 / 1500: loss 8.999750  
iteration 500 / 1500: loss 8.999812  
iteration 600 / 1500: loss 8.999726  
iteration 700 / 1500: loss 8.999796  
iteration 800 / 1500: loss 8.999877  
iteration 900 / 1500: loss 8.999825  
iteration 1000 / 1500: loss 8.999736  
iteration 1100 / 1500: loss 8.999741  
iteration 1200 / 1500: loss 8.999696  
iteration 1300 / 1500: loss 8.999707  
iteration 1400 / 1500: loss 8.999687  
iteration 0 / 1500: loss 172.675160  
iteration 100 / 1500: loss 9.000009  
iteration 200 / 1500: loss 8.999969  
iteration 300 / 1500: loss 9.000000  
iteration 400 / 1500: loss 8.999967  
iteration 500 / 1500: loss 8.999976  
iteration 600 / 1500: loss 8.999971  
iteration 700 / 1500: loss 9.000044  
iteration 800 / 1500: loss 8.999956  
iteration 900 / 1500: loss 8.999927



iteration 1000 / 1500: loss 9.000006  
iteration 1100 / 1500: loss 8.999998  
iteration 1200 / 1500: loss 9.000052  
iteration 1300 / 1500: loss 9.000013  
iteration 1400 / 1500: loss 9.000010  
lr 5.000000e-08 reg 1.000000e+03 train accuracy: 0.098429 val accuracy: 0.098000  
lr 5.000000e-08 reg 5.000000e+03 train accuracy: 0.113735 val accuracy: 0.117000  
lr 5.000000e-08 reg 1.000000e+04 train accuracy: 0.132571 val accuracy: 0.123000  
lr 5.000000e-08 reg 2.000000e+04 train accuracy: 0.177735 val accuracy: 0.192000  
lr 5.000000e-08 reg 3.000000e+04 train accuracy: 0.246490 val accuracy: 0.239000  
lr 5.000000e-08 reg 4.000000e+04 train accuracy: 0.382469 val accuracy: 0.400000  
lr 5.000000e-08 reg 5.000000e+04 train accuracy: 0.416327 val accuracy: 0.408000  
lr 5.000000e-08 reg 1.000000e+05 train accuracy: 0.417878 val accuracy: 0.422000  
lr 1.000000e-07 reg 1.000000e+03 train accuracy: 0.123122 val accuracy: 0.134000  
lr 1.000000e-07 reg 5.000000e+03 train accuracy: 0.151082 val accuracy: 0.147000  
lr 1.000000e-07 reg 1.000000e+04 train accuracy: 0.216490 val accuracy: 0.228000  
lr 1.000000e-07 reg 2.000000e+04 train accuracy: 0.405327 val accuracy: 0.409000  
lr 1.000000e-07 reg 3.000000e+04 train accuracy: 0.417612 val accuracy: 0.419000  
lr 1.000000e-07 reg 4.000000e+04 train accuracy: 0.422959 val accuracy: 0.420000  
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.417020 val accuracy: 0.426000  
lr 1.000000e-07 reg 1.000000e+05 train accuracy: 0.423184 val accuracy: 0.435000  
lr 2.000000e-07 reg 1.000000e+03 train accuracy: 0.146694 val accuracy: 0.144000  
lr 2.000000e-07 reg 5.000000e+03 train accuracy: 0.291020 val accuracy: 0.305000  
lr 2.000000e-07 reg 1.000000e+04 train accuracy: 0.415959 val accuracy: 0.406000  
lr 2.000000e-07 reg 2.000000e+04 train accuracy: 0.417551 val accuracy: 0.428000  
lr 2.000000e-07 reg 3.000000e+04 train accuracy: 0.419061 val accuracy: 0.438000  
lr 2.000000e-07 reg 4.000000e+04 train accuracy: 0.416653 val accuracy: 0.412000  
lr 2.000000e-07 reg 5.000000e+04 train accuracy: 0.419898 val accuracy: 0.420000  
lr 2.000000e-07 reg 1.000000e+05 train accuracy: 0.412796 val accuracy: 0.404000

lr 3.000000e-07	reg 1.000000e+03	train accuracy: 0.205796	val accuracy: 0.216000
lr 3.000000e-07	reg 5.000000e+03	train accuracy: 0.397898	val accuracy: 0.393000
lr 3.000000e-07	reg 1.000000e+04	train accuracy: 0.418143	val accuracy: 0.426000
lr 3.000000e-07	reg 2.000000e+04	train accuracy: 0.419245	val accuracy: 0.421000
lr 3.000000e-07	reg 3.000000e+04	train accuracy: 0.420061	val accuracy: 0.429000
lr 3.000000e-07	reg 4.000000e+04	train accuracy: 0.414653	val accuracy: 0.420000
lr 3.000000e-07	reg 5.000000e+04	train accuracy: 0.412490	val accuracy: 0.413000
lr 3.000000e-07	reg 1.000000e+05	train accuracy: 0.414673	val accuracy: 0.422000
lr 4.000000e-07	reg 1.000000e+03	train accuracy: 0.246531	val accuracy: 0.249000
lr 4.000000e-07	reg 5.000000e+03	train accuracy: 0.417490	val accuracy: 0.422000
lr 4.000000e-07	reg 1.000000e+04	train accuracy: 0.418898	val accuracy: 0.421000
lr 4.000000e-07	reg 2.000000e+04	train accuracy: 0.417347	val accuracy: 0.421000
lr 4.000000e-07	reg 3.000000e+04	train accuracy: 0.416531	val accuracy: 0.410000
lr 4.000000e-07	reg 4.000000e+04	train accuracy: 0.414816	val accuracy: 0.412000
lr 4.000000e-07	reg 5.000000e+04	train accuracy: 0.408816	val accuracy: 0.404000
lr 4.000000e-07	reg 1.000000e+05	train accuracy: 0.420000	val accuracy: 0.415000
lr 5.000000e-07	reg 1.000000e+03	train accuracy: 0.261265	val accuracy: 0.303000
lr 5.000000e-07	reg 5.000000e+03	train accuracy: 0.416816	val accuracy: 0.425000
lr 5.000000e-07	reg 1.000000e+04	train accuracy: 0.420408	val accuracy: 0.432000
lr 5.000000e-07	reg 2.000000e+04	train accuracy: 0.412306	val accuracy: 0.414000
lr 5.000000e-07	reg 3.000000e+04	train accuracy: 0.416204	val accuracy: 0.407000
lr 5.000000e-07	reg 4.000000e+04	train accuracy: 0.417796	val accuracy: 0.417000
lr 5.000000e-07	reg 5.000000e+04	train accuracy: 0.414102	val accuracy: 0.418000
lr 5.000000e-07	reg 1.000000e+05	train accuracy: 0.407510	val accuracy: 0.405000
lr 1.000000e-06	reg 1.000000e+03	train accuracy: 0.408980	val accuracy: 0.401000
lr 1.000000e-06	reg 5.000000e+03	train accuracy: 0.416551	val accuracy: 0.415000
lr 1.000000e-06	reg 1.000000e+04	train accuracy: 0.417694	val accuracy:

```
cy: 0.412000
lr 1.000000e-06 reg 2.000000e+04 train accuracy: 0.415306 val accuracy: 0.412000
lr 1.000000e-06 reg 3.000000e+04 train accuracy: 0.419918 val accuracy: 0.420000
lr 1.000000e-06 reg 4.000000e+04 train accuracy: 0.407918 val accuracy: 0.404000
lr 1.000000e-06 reg 5.000000e+04 train accuracy: 0.415796 val accuracy: 0.411000
lr 1.000000e-06 reg 1.000000e+05 train accuracy: 0.405612 val accuracy: 0.410000
lr 5.000000e-06 reg 1.000000e+03 train accuracy: 0.420694 val accuracy: 0.421000
lr 5.000000e-06 reg 5.000000e+03 train accuracy: 0.411082 val accuracy: 0.416000
lr 5.000000e-06 reg 1.000000e+04 train accuracy: 0.412694 val accuracy: 0.416000
lr 5.000000e-06 reg 2.000000e+04 train accuracy: 0.411633 val accuracy: 0.413000
lr 5.000000e-06 reg 3.000000e+04 train accuracy: 0.398939 val accuracy: 0.408000
lr 5.000000e-06 reg 4.000000e+04 train accuracy: 0.364837 val accuracy: 0.356000
lr 5.000000e-06 reg 5.000000e+04 train accuracy: 0.369918 val accuracy: 0.374000
lr 5.000000e-06 reg 1.000000e+05 train accuracy: 0.325653 val accuracy: 0.325000
best validation accuracy achieved during cross-validation: 0.438000
```

In [14]:

```
# Evaluate your trained SVM on the test set
y_test_pred = best_svm.predict(X_test_feats)
test_accuracy = np.mean(y_test == y_test_pred)
print(test_accuracy)
```

0.426

In [15]:

```
# An important way to gain intuition about how an algorithm works is to
# visualize the mistakes that it makes. In this visualization, we show examples
# of images that are misclassified by our current system. The first column
# shows images that our system labeled as "plane" but whose true label is
# something other than "plane".

examples_per_class = 8
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for cls, cls_name in enumerate(classes):
    idxs = np.where((y_test != cls) & (y_test_pred == cls))[0]
    idxs = np.random.choice(idxs, examples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt.subplot(examples_per_class, len(classes), i * len(classes) + cls + 1)

        plt.imshow(X_test[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls_name)
plt.show()
```



## Inline question 1:

Describe the misclassification results that you see. Do they make sense?

Some seem to make sense - for example, peripheral view of bird that can arguably look like a plane taking off or a dog on a grassy field standing on 4 legs mistaken for a horse. But many other misclassifications are hard to interpret.

# Neural Network on image features

Earlier in this assignment we saw that training a two-layer neural network on raw pixels achieved better classification performance than linear classifiers on raw pixels. In this notebook we have seen that linear classifiers on image features outperform linear classifiers on raw pixels.

For completeness, we should also try training a neural network on image features. This approach should outperform all previous approaches: you should easily be able to achieve over 55% classification accuracy on the test set; our best model achieves about 60% classification accuracy.

In [29]:

```
# Preprocessing: Remove the bias dimension  
# Make sure to run this cell only ONCE  
print(X_train_feats.shape)  
X_train_feats = X_train_feats[:, :-1]  
X_val_feats = X_val_feats[:, :-1]  
X_test_feats = X_test_feats[:, :-1]  
  
print(X_train_feats.shape)
```

```
(49000, 161)
```

```
(49000, 160)
```

In [33]:

```
from cs231n.classifiers.neural_net import TwoLayerNet

input_dim = X_train_feats.shape[1]
hidden_dim = 500
num_classes = 10

#####
# TODO: Train a two-layer neural network on image features. You may want to #
# cross-validate various parameters as in previous sections. Store your best #
# model in the best_net variable. #
#####
results = {}
best_val = -1
best_net = None
best_params = None
learning_rates = [1e-1, 5e-1, 1e0]
regularization_strengths = [1e-3, 5e-3, 1e-2]

for l in learning_rates:
    for r in regularization_strengths:
        print(l)
        net = TwoLayerNet(input_dim, hidden_dim, num_classes)
        curr_loss = net.train(X_train_feats, y_train, X_val_feats, y_val, learning_rate=l, reg=r, num_iters=1500)
        y_train_pred = net.predict(X_train_feats)
        y_train_acc = np.mean(y_train == y_train_pred)
        y_val_pred = net.predict(X_val_feats)
        y_val_acc = np.mean(y_val == y_val_pred)
        results[(l, r)] = (y_train_acc, y_val_acc)
        if y_val_acc > best_val:
            best_net = net
            best_val = y_val_acc
            best_params = (l, r)

#####
#                               END OF YOUR CODE                               #
#####
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' % best_val)
print(best_params)
```

```
0.1
0.1
0.1
0.5
0.5
0.5
1.0
1.0
1.0
lr 1.000000e-01 reg 1.000000e-03 train accuracy: 0.540918 val accuracy: 0.525000
lr 1.000000e-01 reg 5.000000e-03 train accuracy: 0.528388 val accuracy: 0.514000
lr 1.000000e-01 reg 1.000000e-02 train accuracy: 0.521020 val accuracy: 0.520000
lr 5.000000e-01 reg 1.000000e-03 train accuracy: 0.660694 val accuracy: 0.597000
lr 5.000000e-01 reg 5.000000e-03 train accuracy: 0.565143 val accuracy: 0.559000
lr 5.000000e-01 reg 1.000000e-02 train accuracy: 0.514714 val accuracy: 0.500000
lr 1.000000e+00 reg 1.000000e-03 train accuracy: 0.671020 val accuracy: 0.560000
lr 1.000000e+00 reg 5.000000e-03 train accuracy: 0.555163 val accuracy: 0.519000
lr 1.000000e+00 reg 1.000000e-02 train accuracy: 0.509469 val accuracy: 0.486000
best validation accuracy achieved during cross-validation: 0.597000
(0.5, 0.001)
```

In [34]:

```
# Run your best neural net classifier on the test set. You should be able
# to get more than 55% accuracy.
```

```
test_acc = (best_net.predict(X_test_feats) == y_test).mean()
print(test_acc)
```

```
0.577
```